

Exploring Machine Intelligence

Week 3, Convolutional NNs



Motivation for today



-1	-1	-1
-1	8	-1
-1	-1	-1

3x3 conv
→



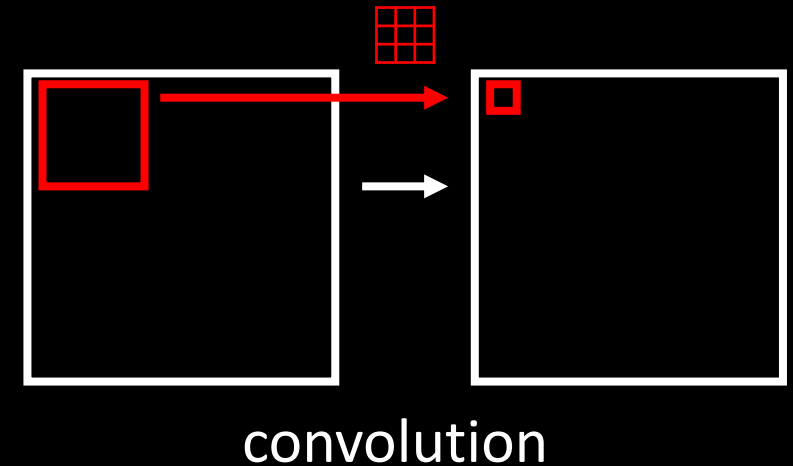
Manual filter: **edge detection**

>>> setosa.io/ev/image-kernels/ <<<

Today


Building a better machine:

- Why should we split data?
- Convolutional operation
- **Convolutional NNs**
- Real world architecture: **AlexNet**
- General feature extractors




Data splitting and overfitting

- Let's illustrate **overfitting** on 2 classes example (dogs vs. cats)



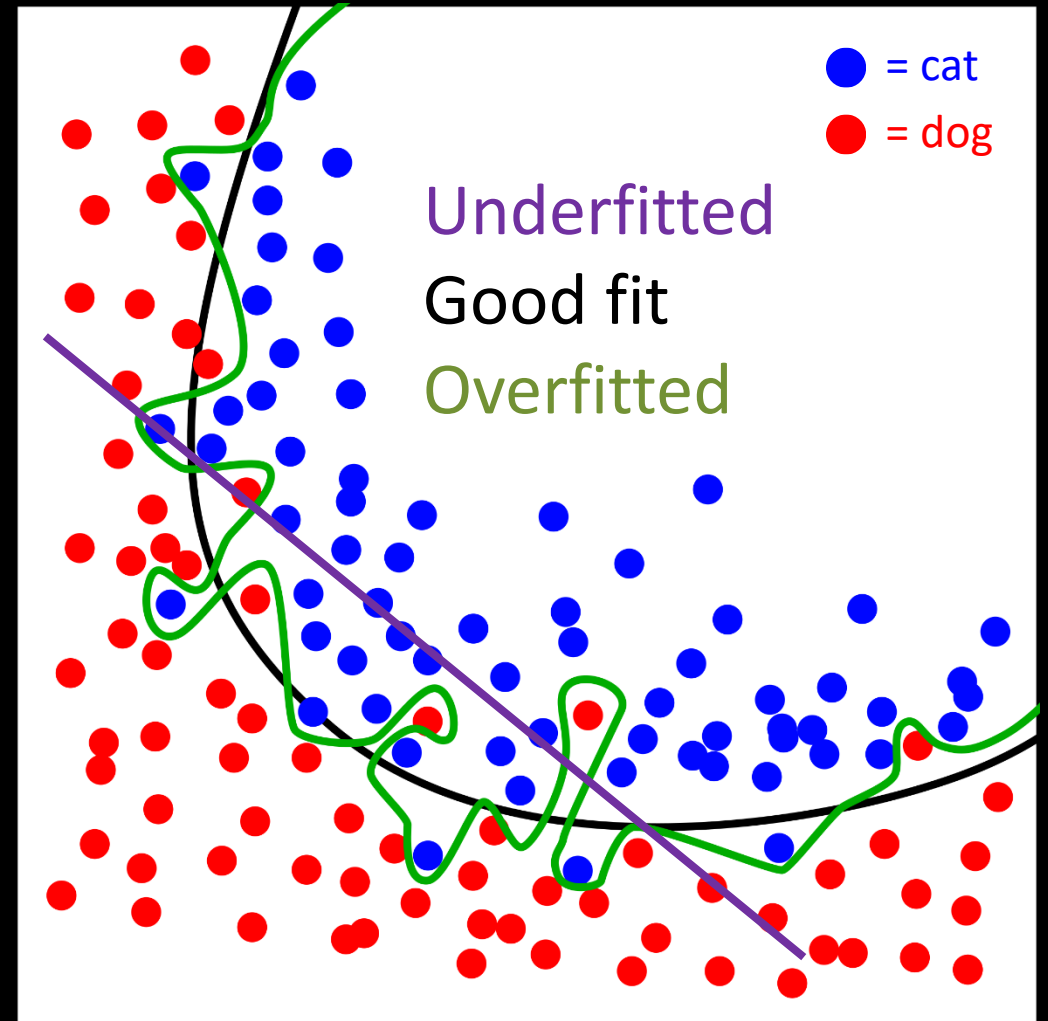
Data splitting and overfitting

- Let's illustrate **overfitting** on 2 classes example (dogs vs. cats)
- We want to split the data with a classifier – model a boundary between the two classes





Data splitting and overfitting

- Let's illustrate **overfitting** on 2 classes example (dogs vs. cats)
- We want to split the data with a classifier – model a boundary between the two classes
- This **separation** can be *sensible* in terms of the boundary it creates ... or it can be *overly complicated*

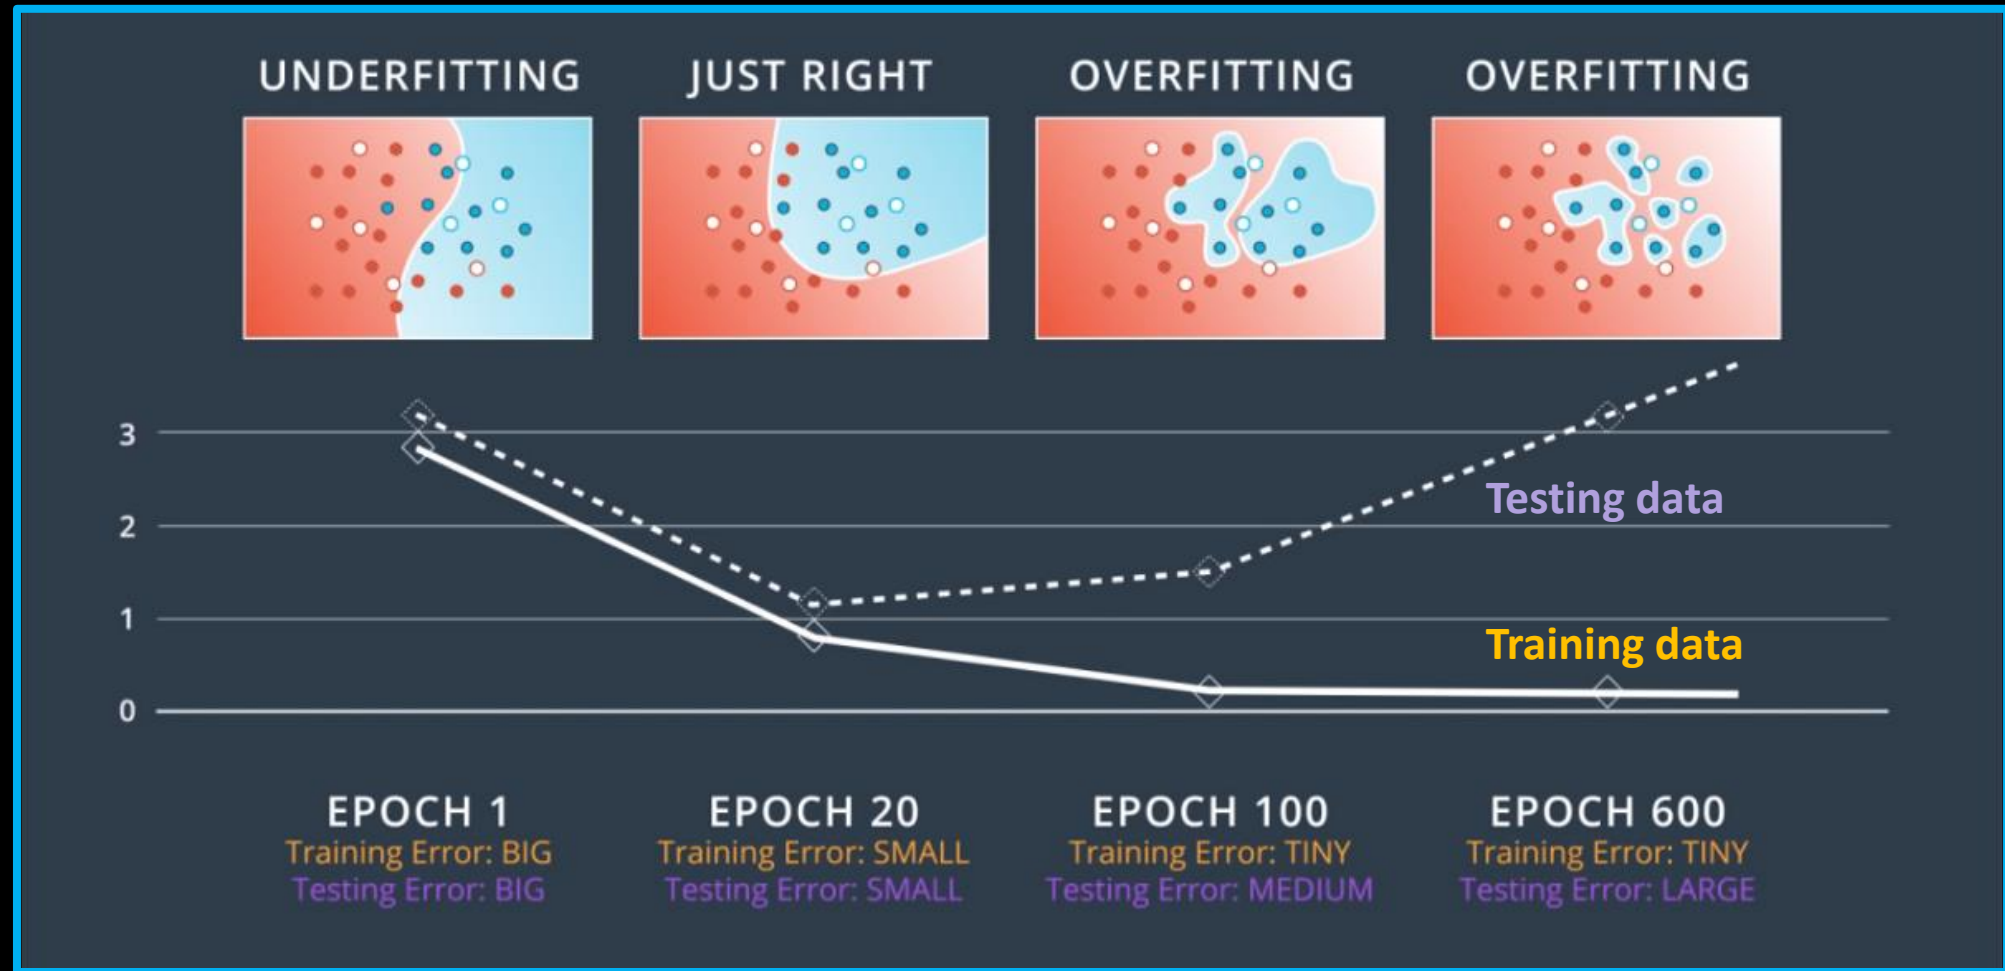


Data splitting and overfitting

- The problem with having the overly **complex boundary** is as follows:
 - It **perfectly separates all the data** we gave it **in the training set!** (100% accuracy on training set)
 - Imagine we have a new samples:  With the boundary being so specialized on this particular training set, the **new samples will fail to be classified correctly.**



- Can we **detect overfitting**? **Yes**:




- Can we **prevent overfitting** from happening?
 - **Smaller models, early stops** in training, various **tricks** ...

Plot from: [here](#)
More intuition: [video](#)

Building a better machine

- To work with images, we need to do better! While something like handwritten digits are reasonably easy to learn with even simple fully connected NNs (around 1998) ... We will need more powerful tools to handle larger datasets (from around 2012).



- One **simplification we made previously**, was that we simply **flattened** the **whole image** and looked at **all pixels independently**.

Images have locality



^^^ *Let's imagine any image in here*

Images have locality



These two pixels will likely contain something relevant to each other.

While these two pixels won't likely have that much in common.

^^^ *Let's imagine any image in here*

Images have locality



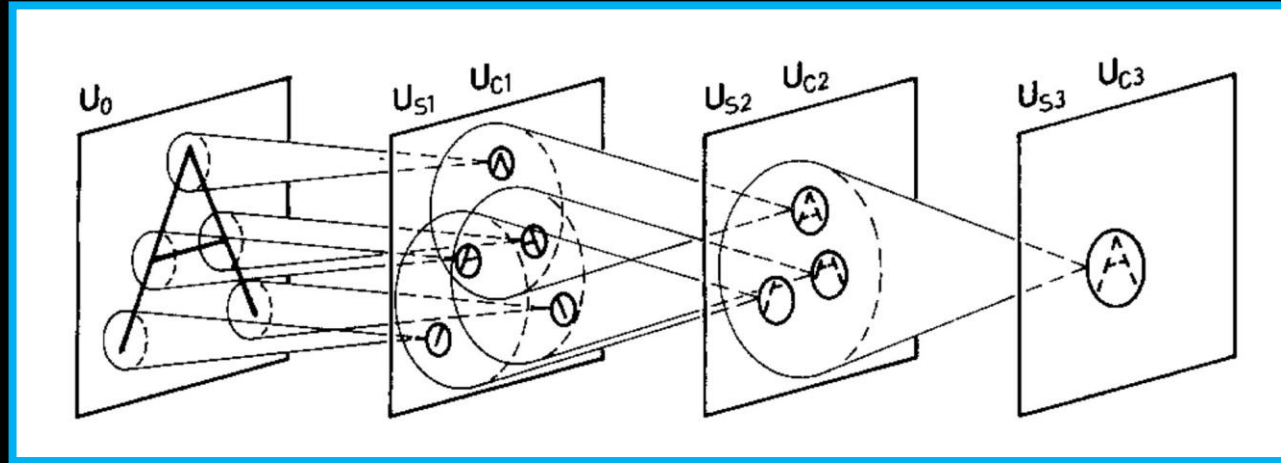
These two pixels will likely contain something relevant to each other.

While these two pixels won't likely have that much in common.

^^^ *Let's imagine any image in here*

- In 1962 a study by *Hubel and Wiesel* explored human visual system and discovered two types of cells – **localized cells** processing details and **complex cells** covering larger areas.

Neocognitron




<- Predecessor to
Convolutional
Neural Networks

- In the 1982 work by Kunihiro Fukushima, they attempted to build an **artificial visual system** which had two types of units:
 - “*s-cells*” focusing on single image features
 - “*c-cells*” combining information from an area

Bonus: [video](#)


Convolution



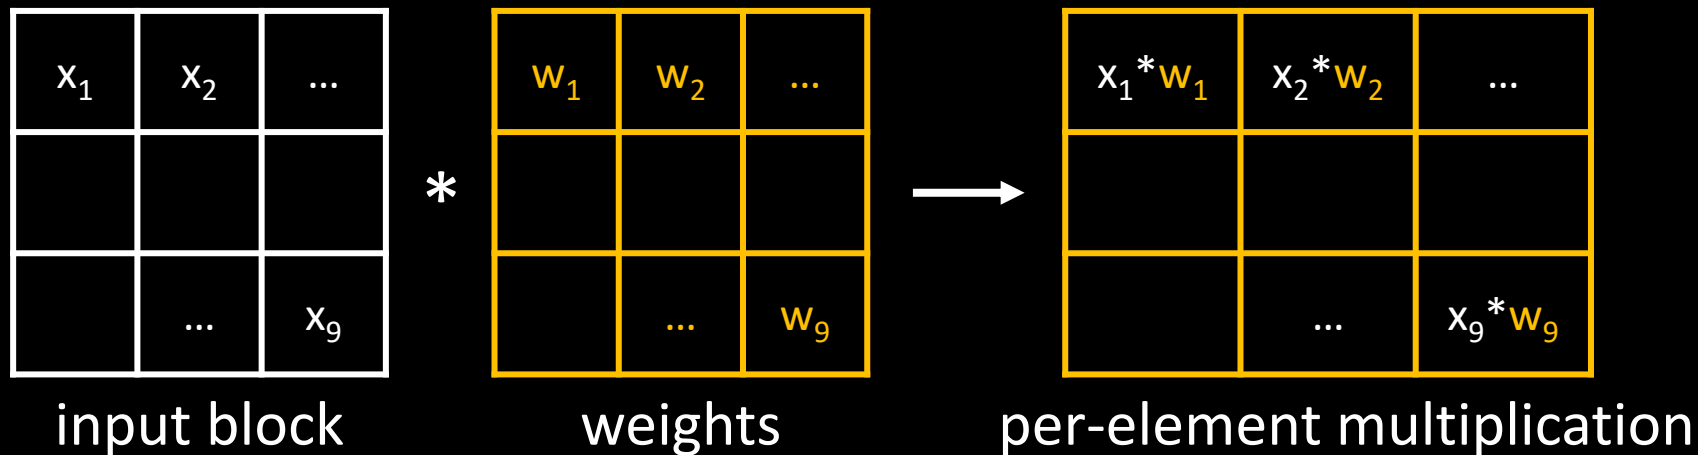
- **Convolution** is an operation applied on an image ...

Convolution

Image:




- **Convolution** is an operation applied on an image ...

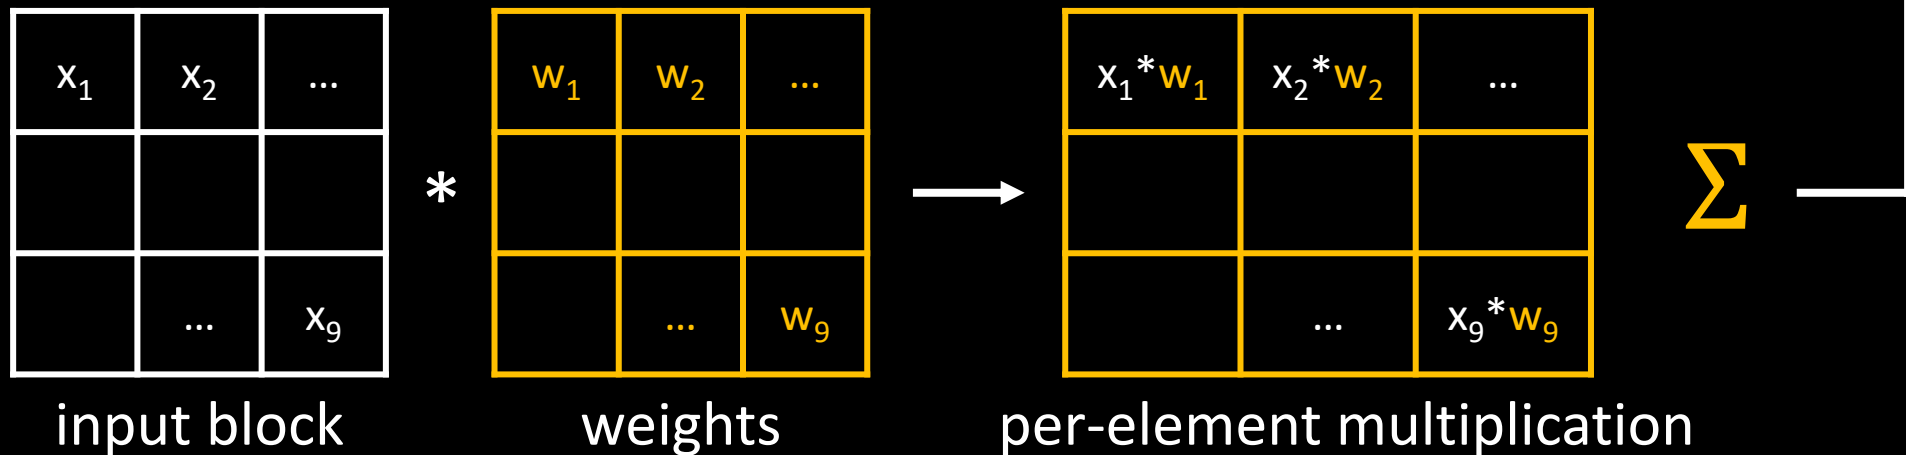


Convolution

Image:




- **Convolution** is an operation applied on an image ...




Convolution

Image:




Convolution

Image:



Convolution

Image:



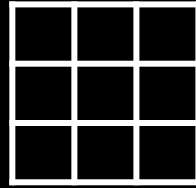
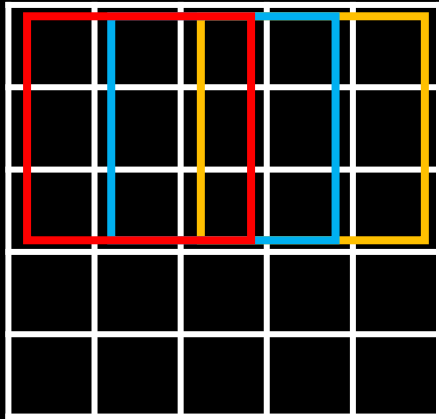
Jump by one
space in input
data

The same
9 weights!


Also jump by
one space in
the outputs

Convolution

Image:




3x3 convolution



Convolution
output: also
an image

- The **same weights** are applied as a **filter**, jumping over the whole image ... effectively producing an image on the output!
- Convolutions were used before neural networks as transformation functions to process images.

Convolution



- **Terminology:** we called these 3x3 numbers “**weights**”, but with convolutions they are also referred to as **kernel**
- **Detail:** To keep the same size of the output image, we need to extend the original image by 1 pixel – we pretend that it includes zeros (also called “**zero padding**”)

Manual convolution filters



1/16	1/8	1/16
1/8	1/4	1/8
1/16	1/8	1/16

3x3 conv
→



blur


>>> setosa.io/ev/image-kernels/ <<<

Manual convolution filters



-2	-1	0
-1	1	1
0	1	2

3x3 conv



emboss

>>> setosa.io/ev/image-kernels/ <<<

Manual convolution filters



-1	-1	-1
-1	8	-1
-1	-1	-1

3x3 conv



outline

>>> setosa.io/ev/image-kernels/ <<<


Automatic convolution filters?

- Similarly as in the previous cases (and your homework), you can manually find values for weights in Neural Network models to process the given dataset in some way (minimizing error).
- Or you can try to get these **parameters of convolution kernels automatically** (like any other parameters of the model) **during training**.

Neural Networks with Convolution

- In addition to the **fully connected layers** of neurons:

Fully connected / Dense
layer




Parameters in
connections

Neural Networks with Convolution

- In addition to the **fully connected layers** of neurons:


Fully connected / Dense layer



Parameters in connections

- We can also use the **convolutional layers**:


Convolutional layer



Parameters in the kernel of the convolution operation


Difference in processing images

Fully connected / Dense layer



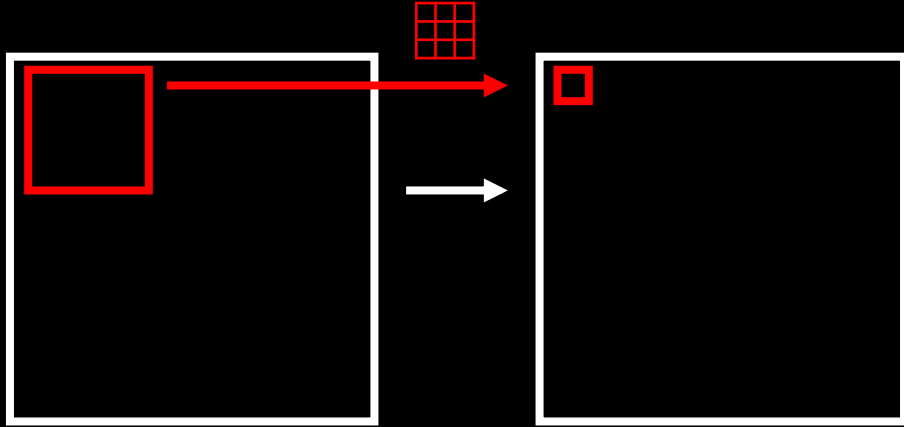
of Weights = # of Pixels
Multiplied by number of
neurons in layer.
... a lot!

Convolutional layer




**# of Weights = size of
kernel**
... fixed for any input size!

Convolutions and Pooling




convolution

Convolutions and Pooling



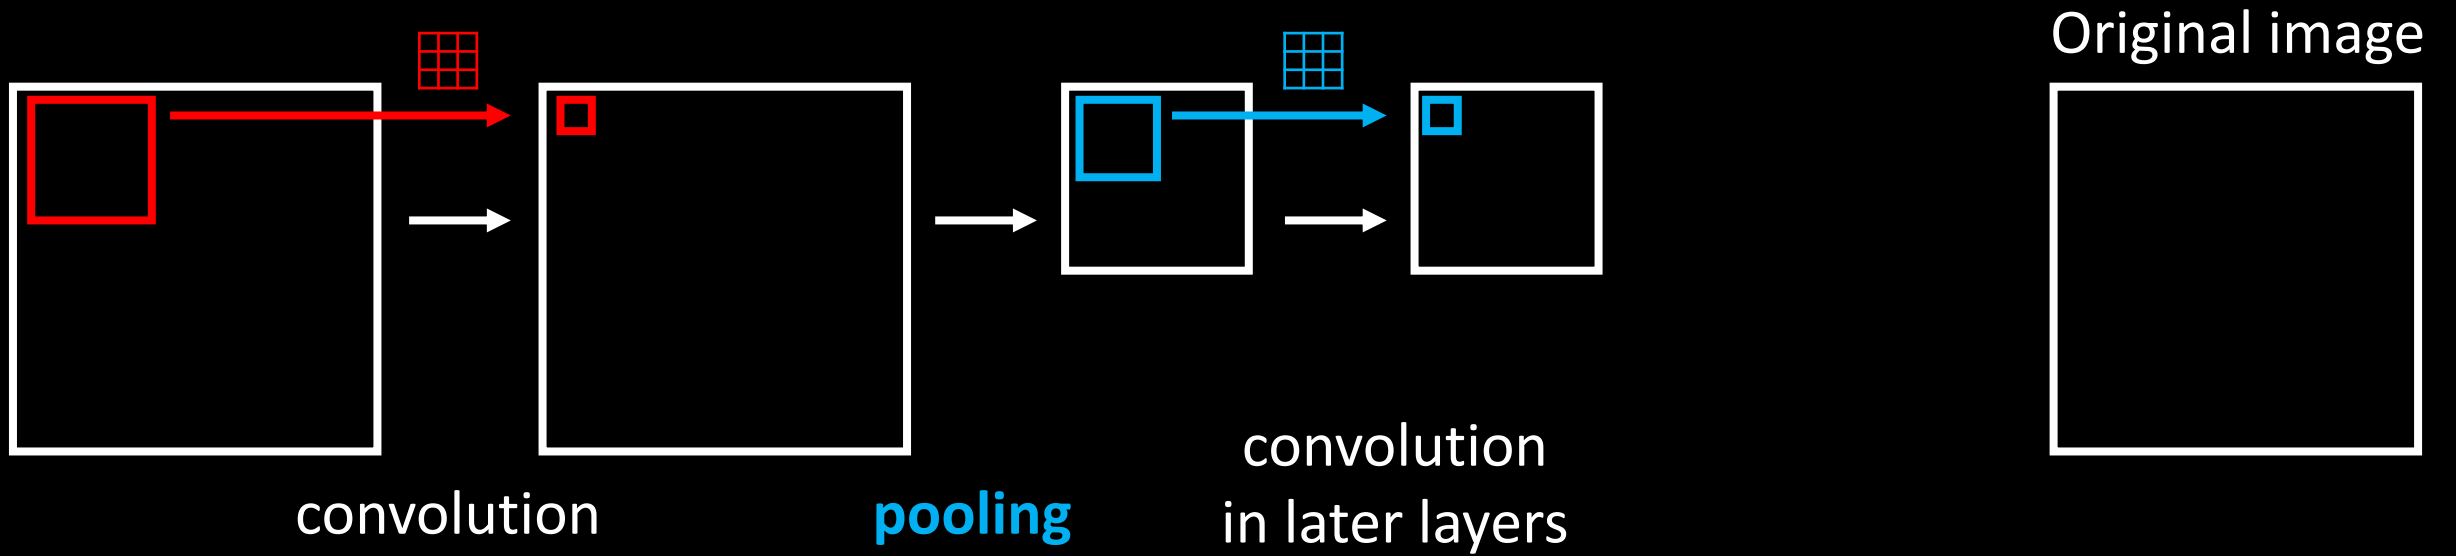
- Convolutional layers are usually followed by **downsampling layers (pooling)** which rescale the image (less information)

Convolutions and Pooling



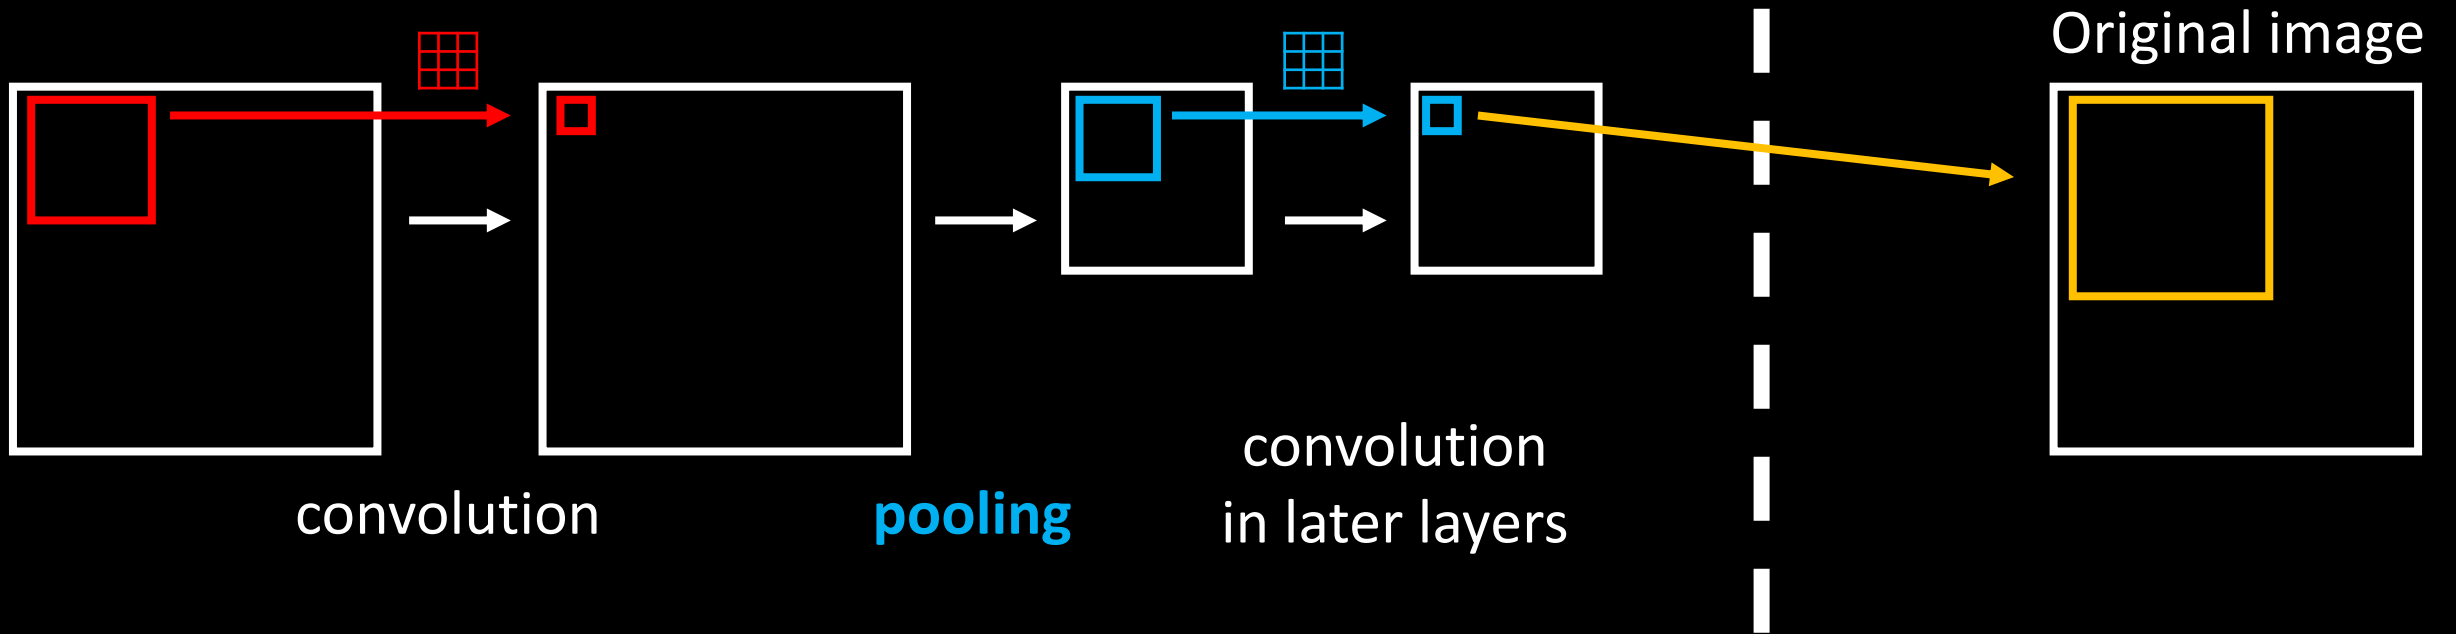
- Convolutional layers are usually followed by **downsampling layers (pooling)** which rescale the image (**less information**)

Convolutions and Pooling



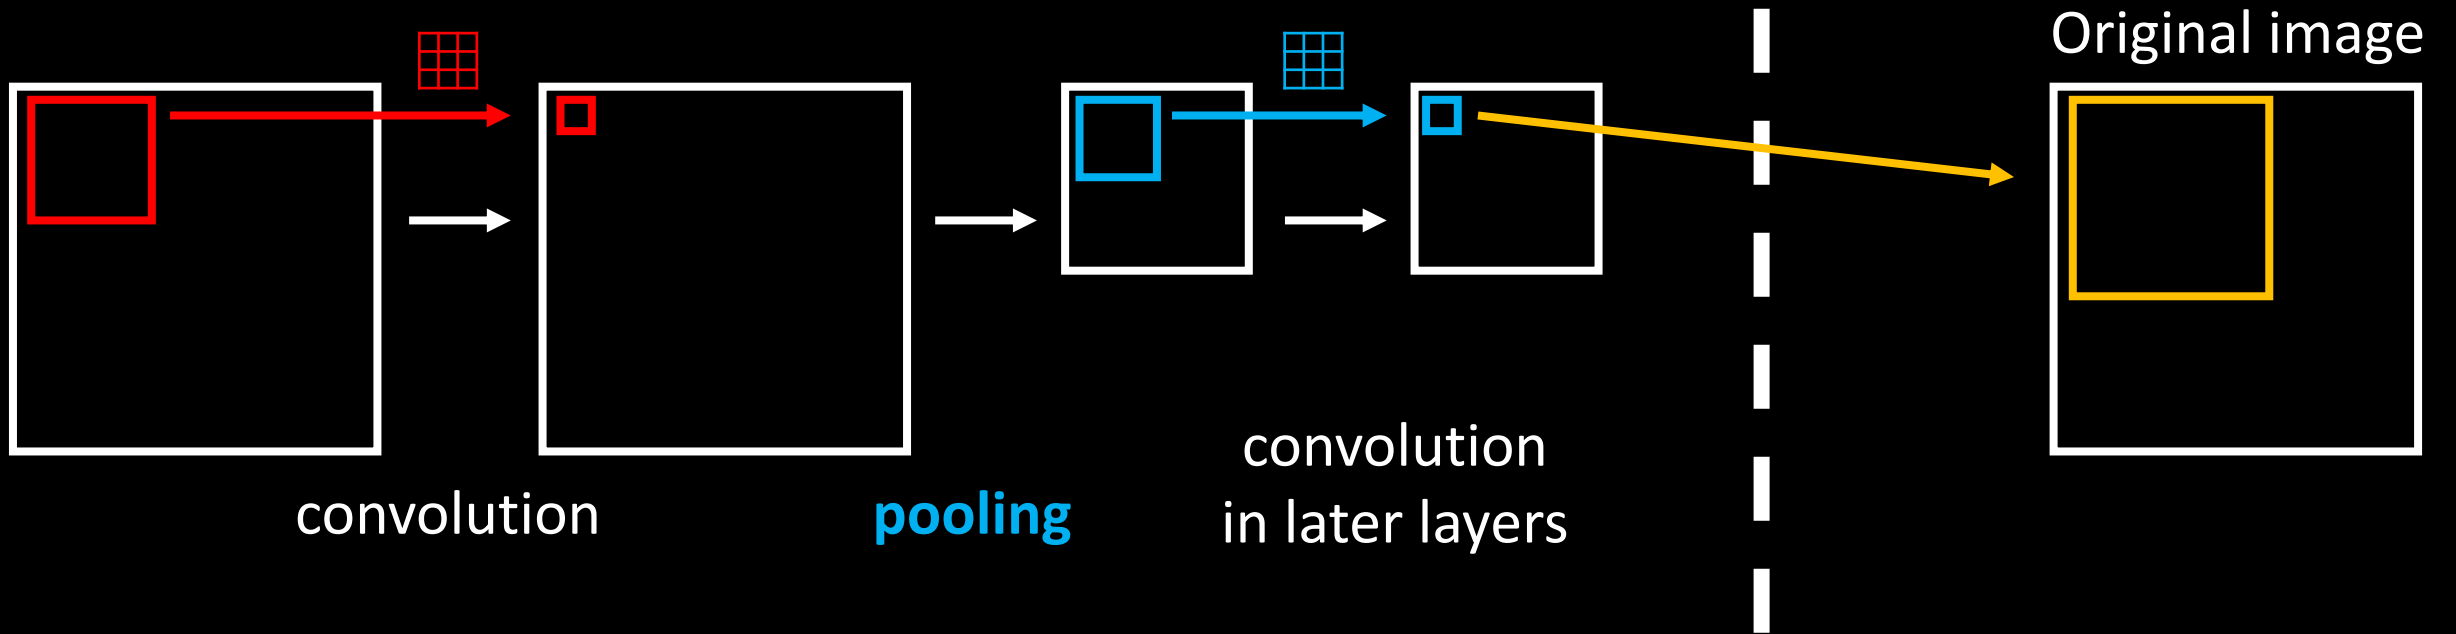
- Convolutional layers are usually followed by **downsampling layers (pooling)** which rescale the image (less information)
- “Later” convolutions are also looking at much **larger region in the original image** with their kernel

Convolutions and Pooling



- Convolutional layers are usually followed by **downsampling layers (pooling)** which rescale the image (less information)
- “Later” convolutions are also looking at much **larger region in the original image** with their kernel

Convolutions and Pooling



- This means, that the later convolutional layers can **specialize on processing different scales of the image** (This is huge!)



details

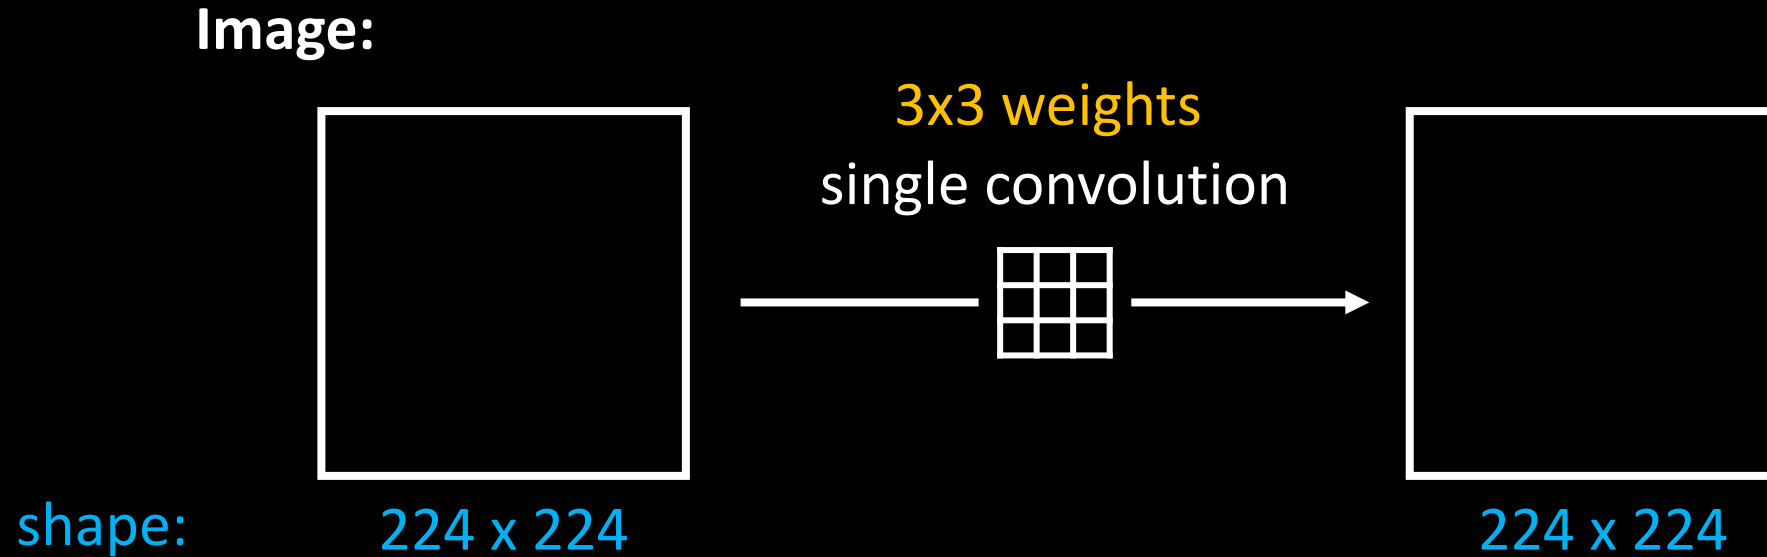


larger areas

... We will see that this means visually!

*One more detail
before the pause*

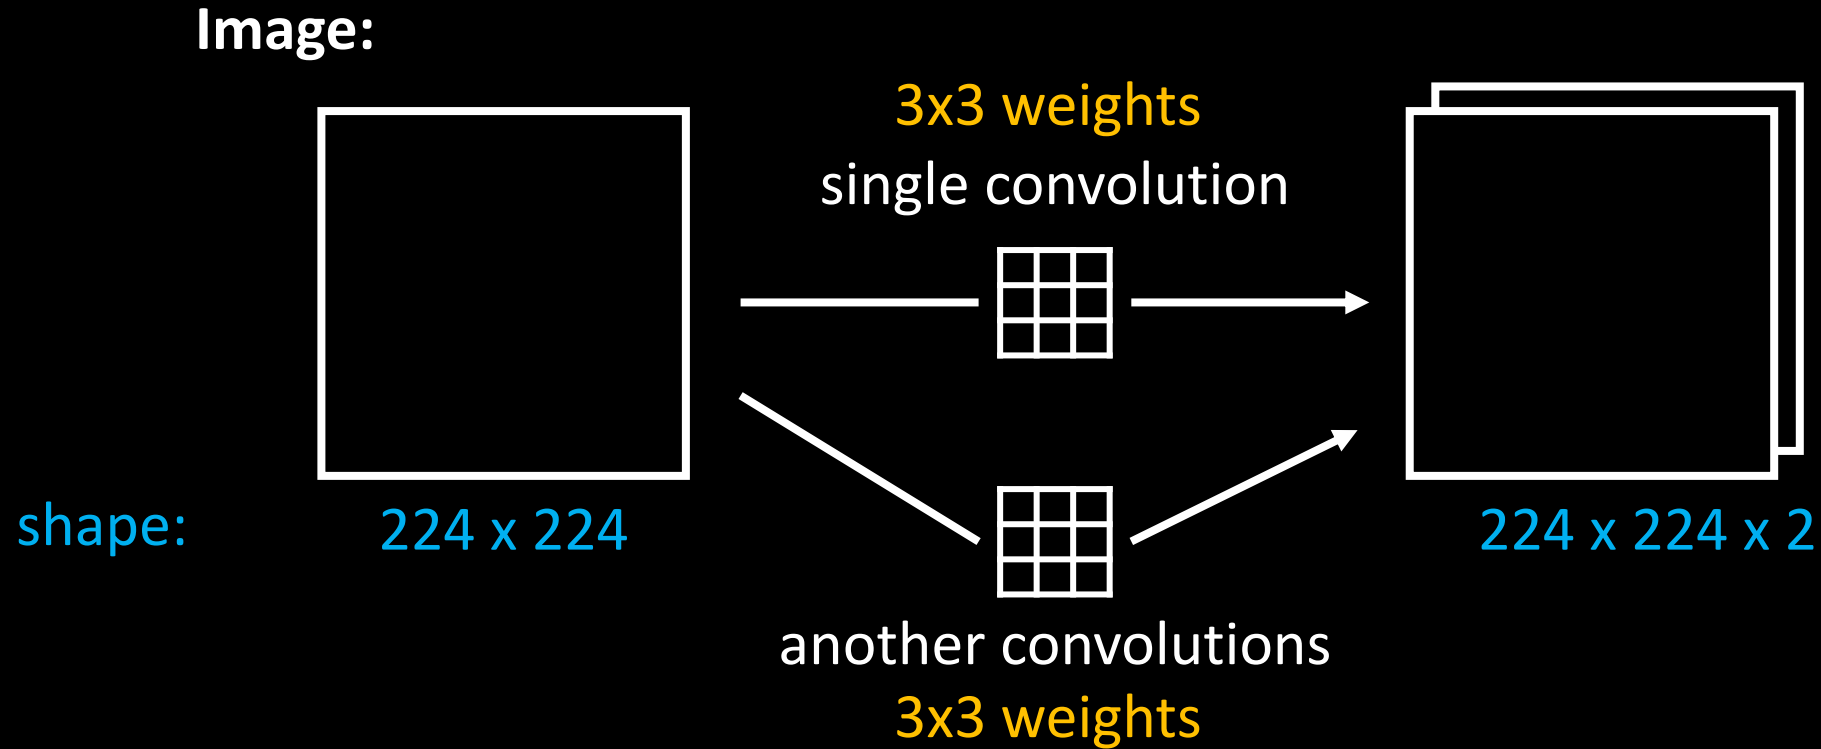
Shapes and weights!



- It can be slightly confusing to read what happens in NN architecture graphs – this should help ^^

*One more detail
before the pause*

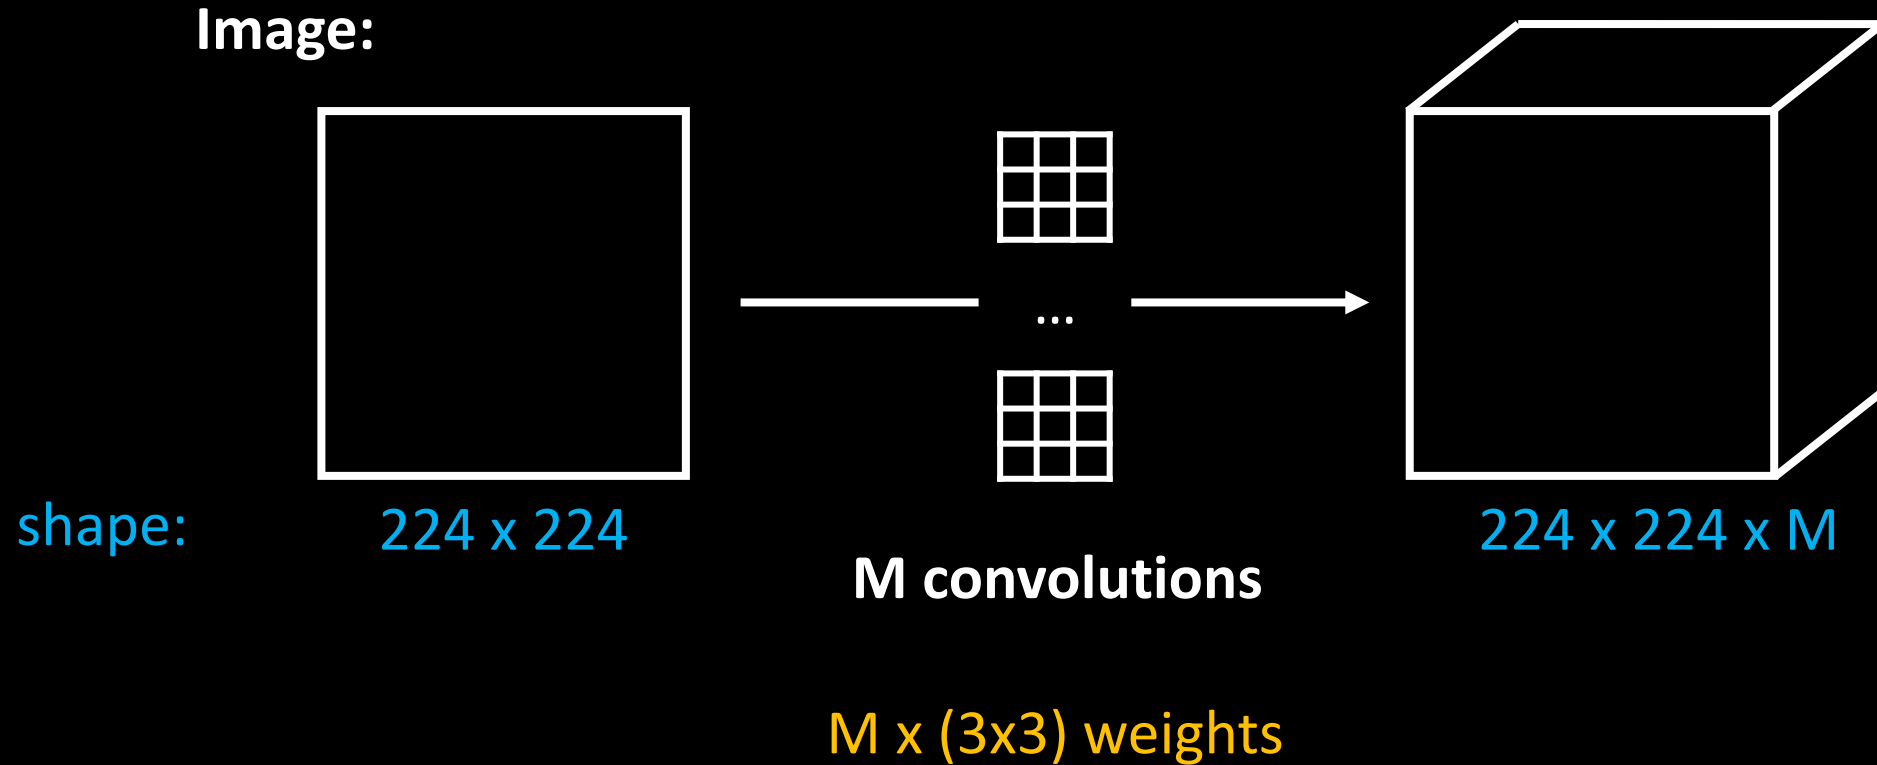
Shapes and weights!



- It can be slightly confusing to read what happens in NN architecture graphs – this should help ^^

*One more detail
before the pause*

Shapes and weights!

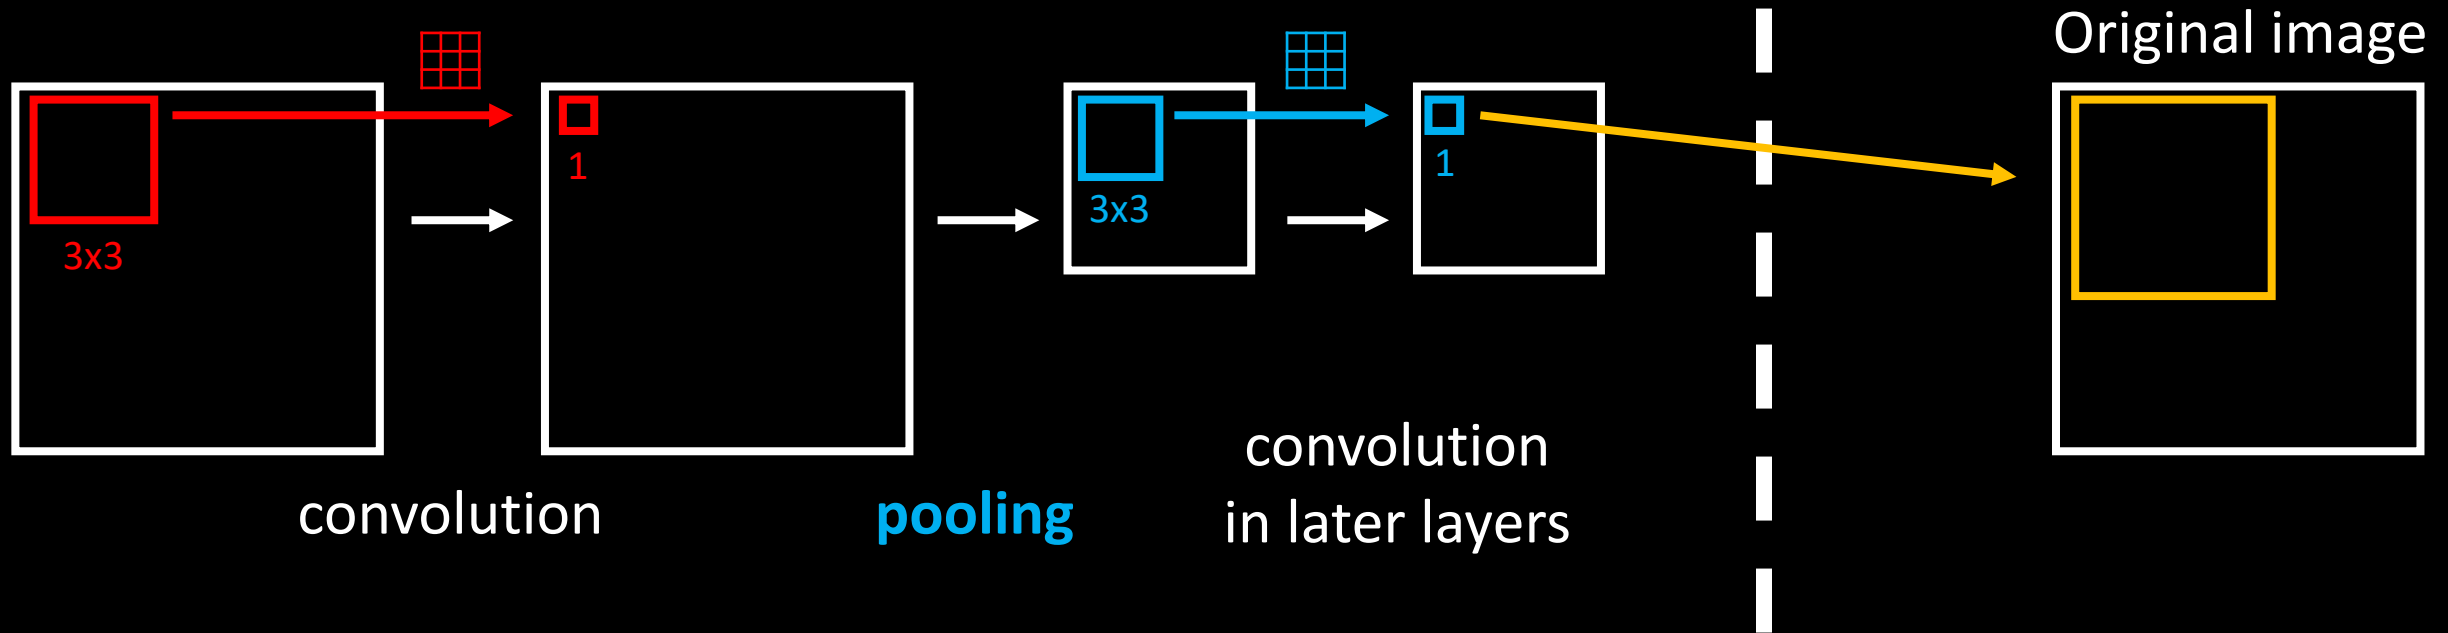


- It can be slightly confusing to read what happens in NN architecture graphs – this should help ^^

Pause 1

Specialized filters

- The (**Convolution -> Pooling**)^{REPEAT} **combo**:



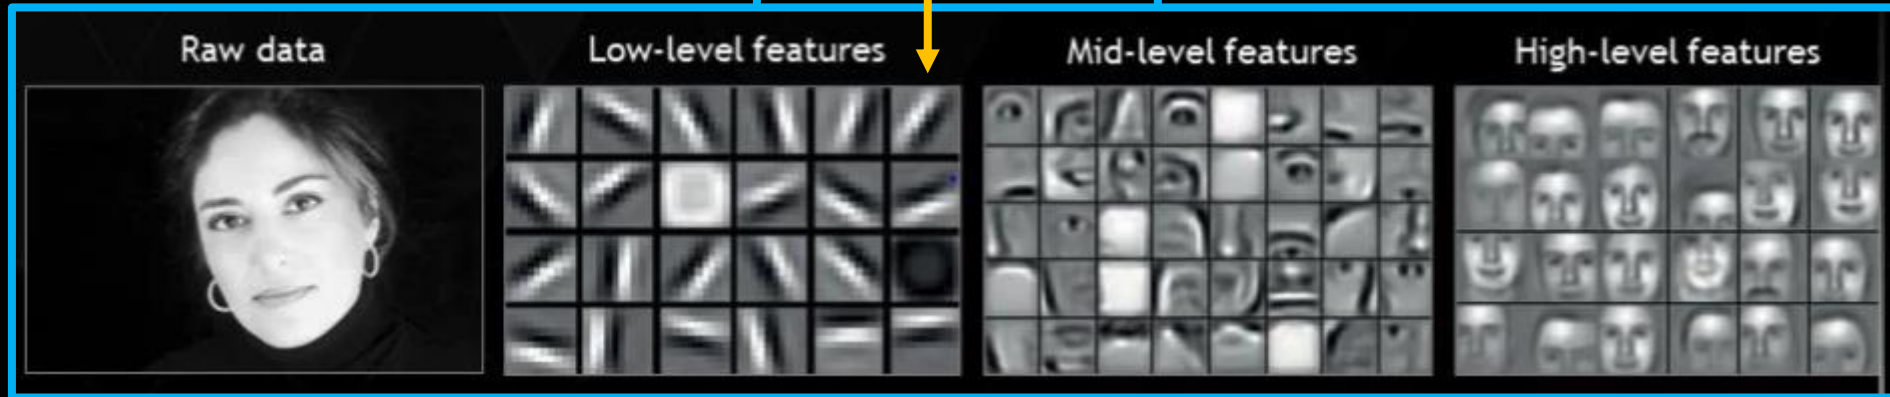
- Allows convolutional layers to **specialize on different scales!**

Specialized filters

- Local features
- Shapes, colors, edges, ...

kernel
visualizations

- Higher level features, more general
- Objects, concepts ...



Earlier layers

depth

Later layers

Intuition for classification


Classifications

High level
features

Mid level
features

Low level
features

Inputs



Bonus: [novel methods of visualizations](#)

Intuition for generation

Let's say we can hack into Convolutional layers of generative models.

- If we target convolutions of **high level features**: *large concepts*



Intuition for generation

Let's say we can hack into Convolutional layers of generative models.

- If we target convolutions of **high level features**: *large concepts*



- If we target convolutions of **low level features**: *details, textures*



General Convolutional NN architecture:

- Most of CNN architectures follow this simple layout:


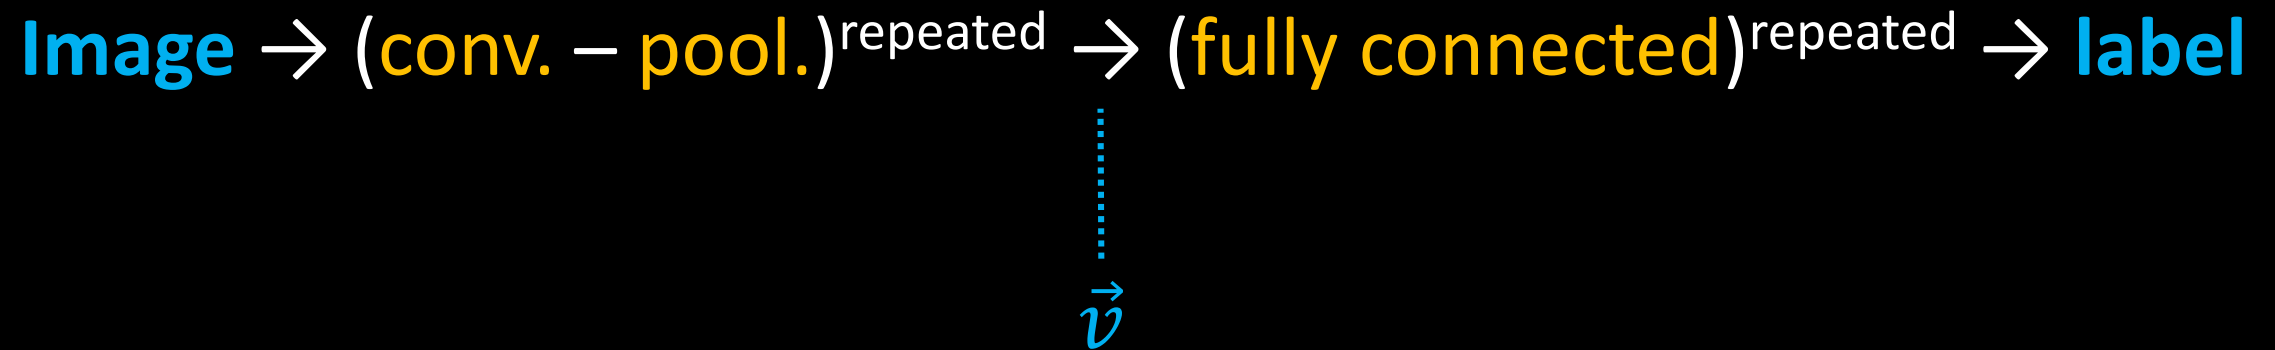


Image \rightarrow (**conv.** – **pool.**)^{repeated} \rightarrow (**fully connected**)^{repeated} \rightarrow **label**

General Convolutional NN architecture:



General Convolutional NN architecture:

Image \rightarrow (**conv.** — **pool.**)^{repeated} \rightarrow (**fully connected**)^{repeated} \rightarrow **label**

Feature extractor

- This part of the model is responsible for **extracting the relevant information** from the image and compressing it into a single vector (the flattened \vec{v})

⋮
 \vec{v}

General Convolutional NN architecture:

Image \rightarrow (conv. — pool.)^{repeated} \rightarrow (fully connected)^{repeated} \rightarrow **label**

Feature extractor

- This part of the model is responsible for **extracting the relevant information** from the image and compressing it into a single vector (the flattened \vec{v})

\vdots
 \vec{v}

Classification

- This part is responsible for **classifying** the vector representation of the image and assigning it a label.

General Convolutional NN architecture:

Image \rightarrow (**conv.** – **pool.**)^{repeated} \rightarrow (**fully connected**)^{repeated} \rightarrow **label**

Feature extractor

- This part of the model is responsible for **extracting the relevant information** from the image and compressing it into a single vector (the flattened \vec{v})

Classification

- This part is responsible for **classifying** the vector representation of the image and assigning it a label.

\vec{v}

Feature vector

- A compressed representation of the original data.

General Convolutional NN architecture:

Image \rightarrow (**conv. — pool.**)^{repeated} \rightarrow (**fully connected**)^{repeated} \rightarrow **label**

Feature extractor

- This part of the model is responsible for **extracting the relevant information** from the image and compressing it into a single vector (the flattened \vec{v})

\vdots
 \vec{v}


Feature

Classification

- This part is responsible for **classifying** the vector representation of the image and assigning it a label.

This **differentiation of function emerges automatically** from the architecture design. It is not imposed manually by how we feed the data. (**Data driven** feature extraction and classification is cool!)


Real world example: AlexNet




By now, we know *most of* the layers used in this CNN architecture!

One more detail: the **Dropout layers**

Dropout layers



(a) Standard Neural Net



(b) After applying dropout.

Dropout deletes connections (with chosen probability) during the training of a NN.

Why?

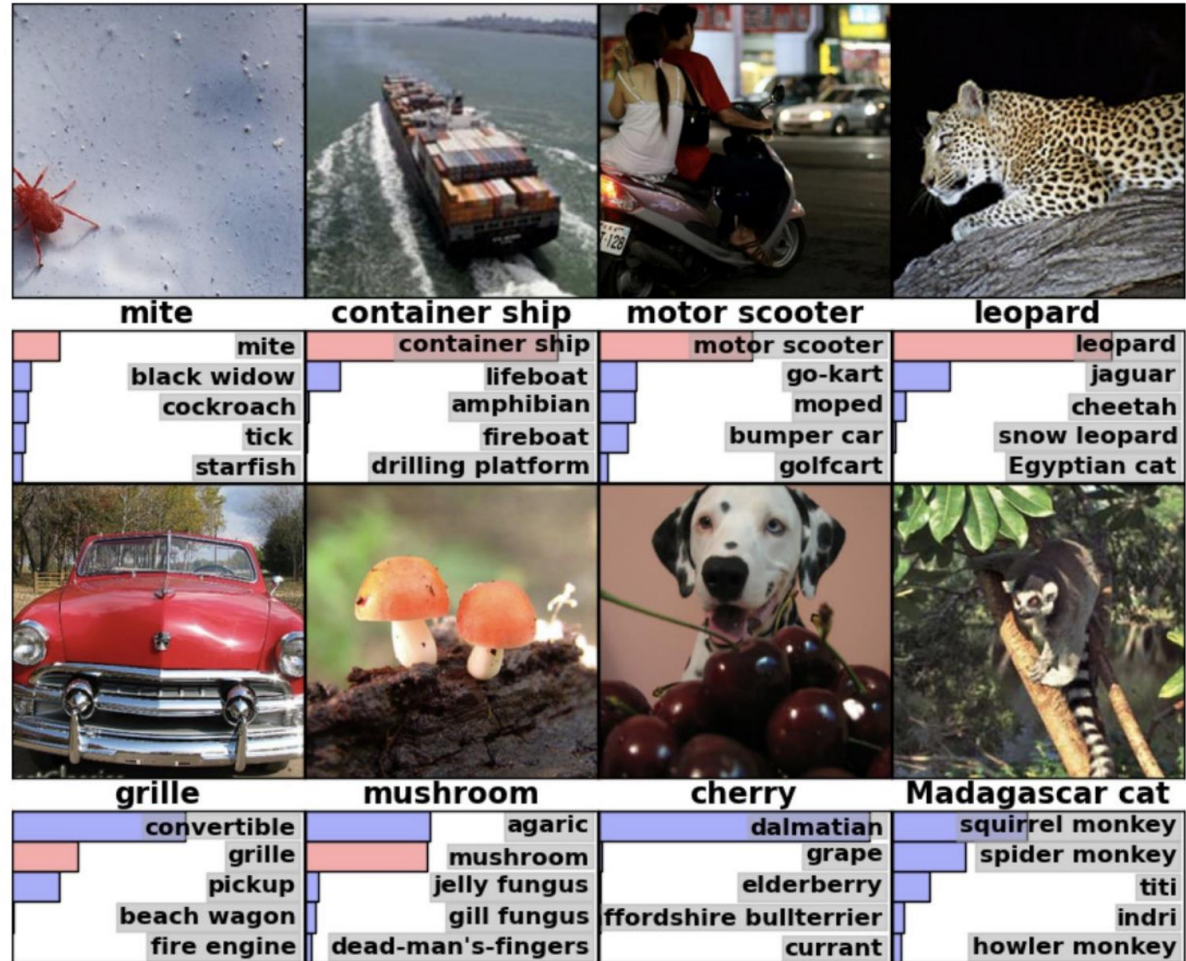
This prevents the network to create overly complicated mappings.

- Dropout is a widely used technique to **helps with overfitting** of NNs.

ImageNet competition


IMAGENET

- 1,000 object classes (categories).
- Images:
 - 1.2 M train
 - 100k test.



Context: the ImageNet competition


Classification: ImageNet Challenge top-5 error



Top-5 error: each model gives top 5 most probable predictions. We count it if the correct class is at least in one of them.

BTW1: Human 5.0%

Context: the ImageNet competition



Top-5 error: each model gives top 5 most probable predictions. We count it if the correct class is at least in one of them.


BTW1: Human 5.0%

BTW2: best today +/- 2.3%

- Since **AlexNet**, all the following models used deep convolutional NNs.

General feature description

Image → Feature extractor → \vec{v} → Classification → **label**



224x224x3 px




vector of 4096 real numbers



1000 classes

General feature description

Image \rightarrow Feature extractor $\rightarrow \vec{v}$



224x224x3 px




vector of 4096 real numbers

- This part of the model becomes useful as a general feature extractor (its learned **representation-ability is data driven** – depends on the dataset!)

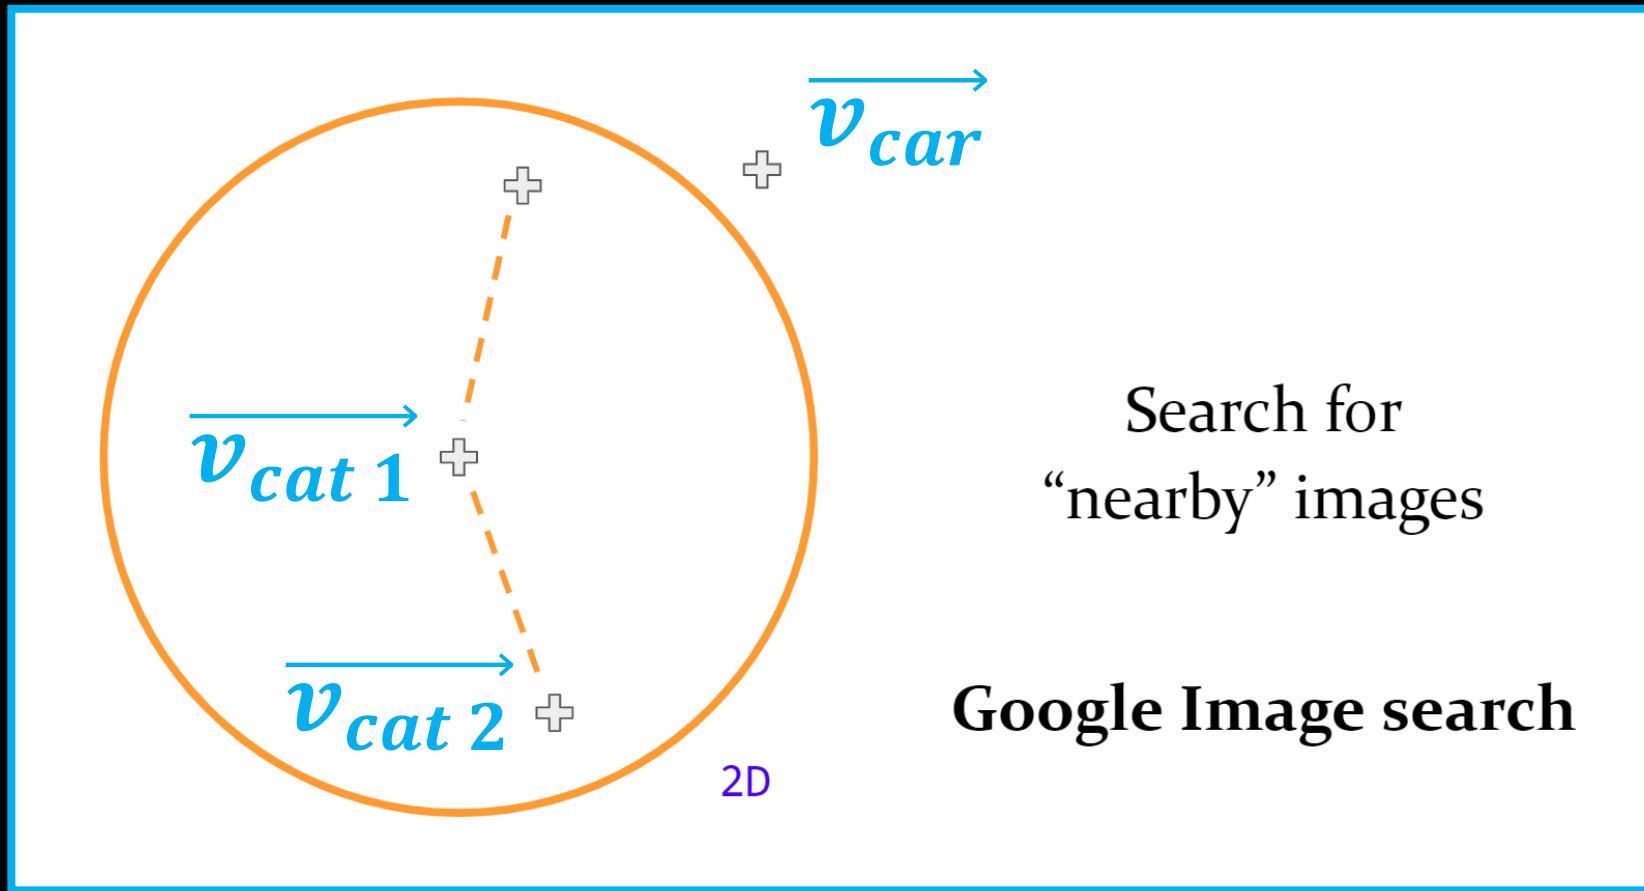
General feature description

Image \rightarrow Feature extractor $\rightarrow \vec{v} \rightarrow$ Classification \rightarrow label




General feature description

Image → **Feature extractor** → \vec{v} → **Classification** → **label**



General feature description

Image → Feature extractor → \vec{v} → Classification → label





Pieter Bruegel the Elder,
ca. 1568
The Tower of Babel
Museum Boijmans Van Beuningen



William Mulready,
1810
Lock gate
Yale Center for British



Unknown, 1828 to 1829
Landscape with Lake
Yale Center for British Art



Orest Dubay
At Sunset
Galéria umelcov
Spiša




Yoo, Hyun-mi
Still Life
Gyeonggi Museum of
Modern Art



Ralph Balson, 1950
Abstraction
Art Gallery of New South Wales



Unknown
Em vermelho
The Adolpho Leimer Collection of
Brazilian Constructive Art at The
Museum of Fine Arts, Houston



Piet Mondrian
Composition with Red, Blue
and Yellow
Kunsthau Zürich

Mario Klingemann - X Degrees of Separation (2018)

Image 1 → Feature extractor → \vec{v}_1

Vectors in the path ← Feature extractor ← Other images

Image 2 → Feature extractor → \vec{v}_2



Pieter Bruegel the Elder,
ca. 1568
The Tower of Babel
Museum Boijmans Van Beuningen



William Mulready,
1810
Lock gate
Yale Center for British



Unknown, 1828 to 1829
Landscape with Lake
Yale Center for British Art



Orest Dubay
At Sunset
Galéria umelcov
Spiša




Yoo, Hyun-mi
Still Life
Gyeonggi Museum of
Modern Art



Ralph Balson, 1950
Abstraction
Art Gallery of New South Wales



Unknown
Em vermelho
The Adolpho Leimer Collection of
Brazilian Constructive Art at The
Museum of Fine Arts, Houston



Piet Mondrian
Composition with Red, Blue
and Yellow
Kunsthaus Zürich

Mario Klingemann - X Degrees of Separation (2018)

Practicum: Convolutional NNs

Continue with code on our Github:


- github repo: github.com/previtus/ci_exploring_machine_intelligence
- notebook: [ml03_cnn_image_search.ipynb](#)

Next class

8.5. Bank holiday = no class!

15.5. Invited guest lecture by Rebecca Fiebrink

- Interaction with Machine Learning models using Wekinator
- Focus on Interaction Design!



Homework reading:

Deadline
8.5.

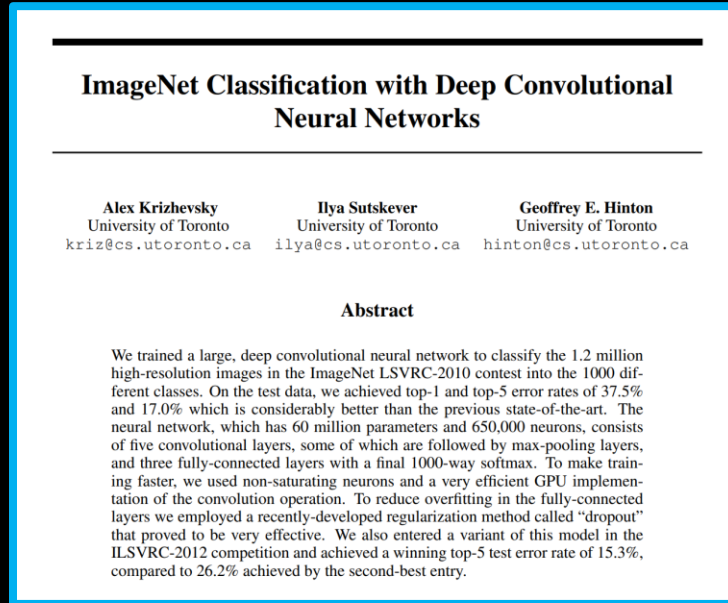
Please read:

Everything, however, you can briefly skim through:

3.2 Training on Multiple GPUs

3.3 Local Response Normalization

3.4 Overlapping Pooling



Tasks:

- In your words, **summarize the paper**.
- **Answer few** prepared **questions**.

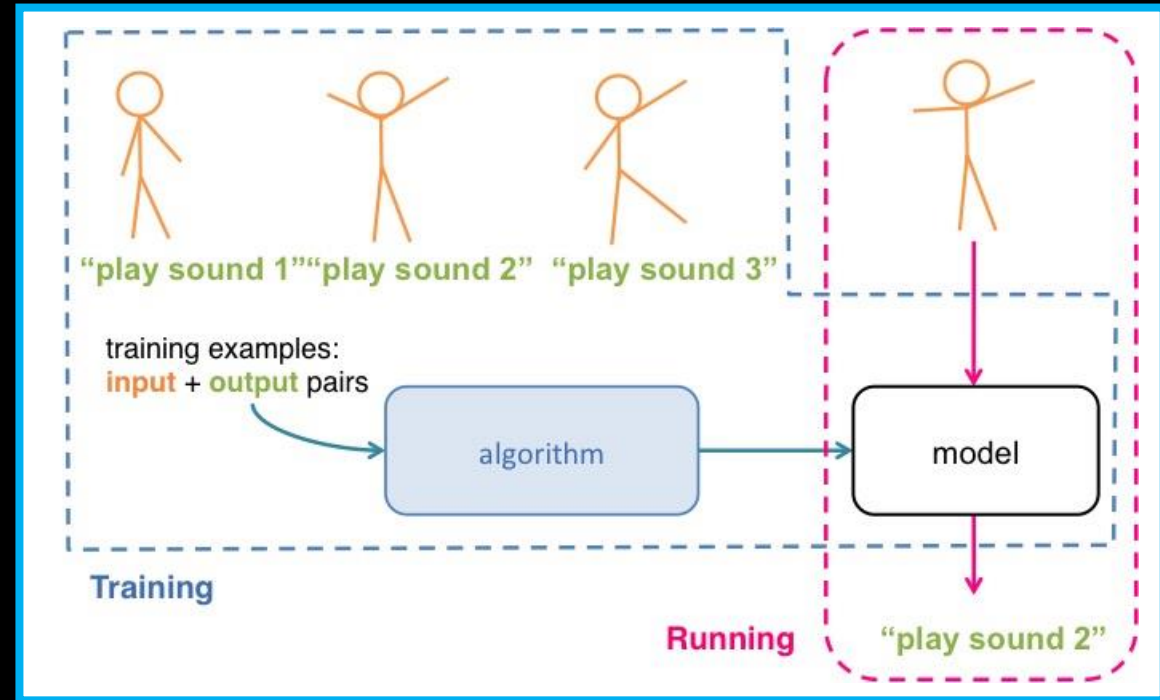
ImageNet Classification with Deep Convolutional Neural Networks
– **A. Krizhevsky**, (2012)

PDF link: papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf

Pre-class preparation:

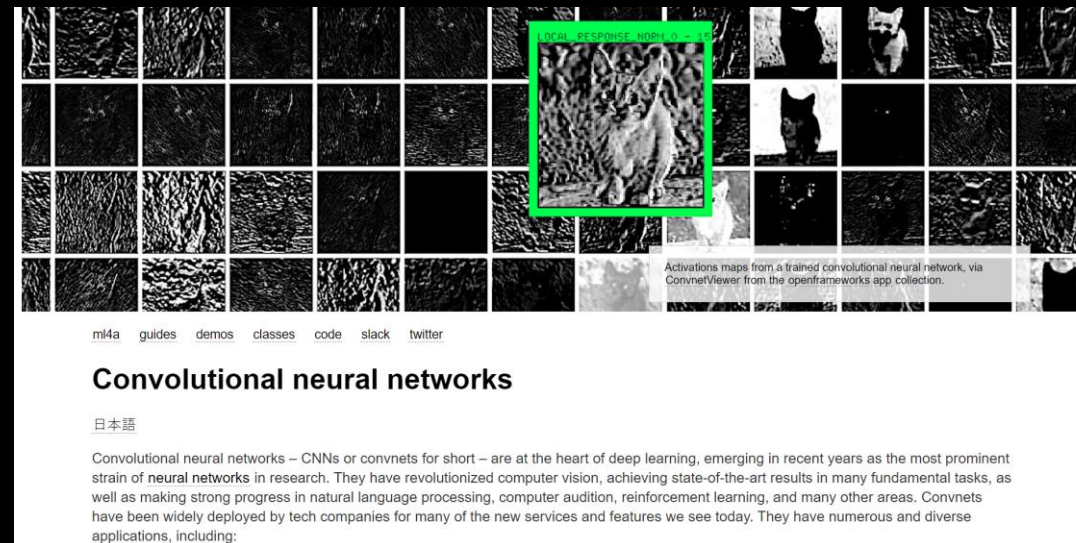
Deadline
15.5.

Get to know the basics
of using **Wekinator**:



Start with the **walkthrough** at: www.wekinator.org/walkthrough/

Bonus links for further reading:



Convnets on ML4A: ml4a.github.io/ml4a/convnets/

Handy notebooks: ml4a.github.io/guides/

The end