

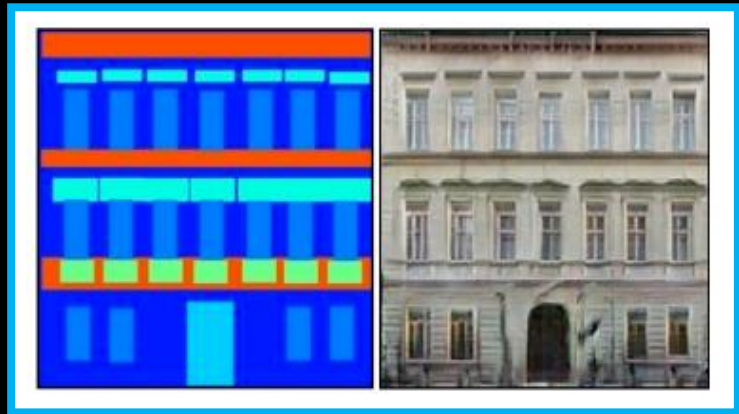
Exploring Machine Intelligence

Week 6, Generative Models II



Motivation for today

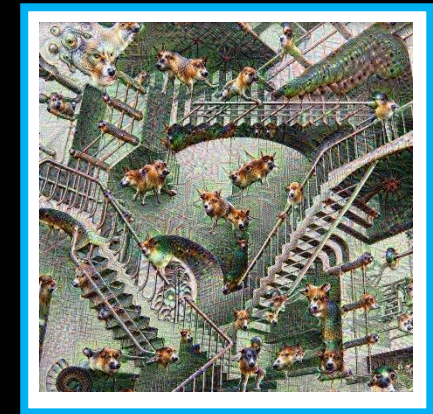
Overview of additional machine learning techniques:



pix2pix



style transfer



deep dream

Today


Overview of some additional Generative Models:

- **Pix2pix** and domain to domain transfer
- **Style transfer** technique
- **Deep Dream** technique


Larger focus on the practical session:

- Using **Progressively Growing GAN** – detailed instructions from data processing to model training

Domain to Domain


- We have seen few models working with images on outputs and inputs (for example [AutoEncoders](#)). Similar architectures can be also used to model **relation between two domains** of data. 

Domain to Domain

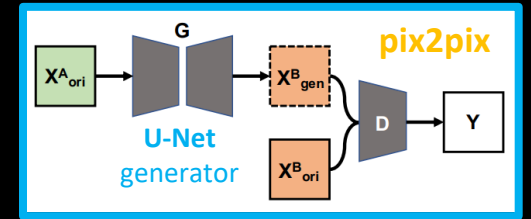
- We have seen few models working with images on outputs and inputs (for example [AutoEncoders](#)). Similar architectures can be also used to model **relation between two domains** of data. 

- Let's look at a predecessor of this idea, the **U-Net model**:
 - Paper with U-Net applied on the task of lunar crater identification

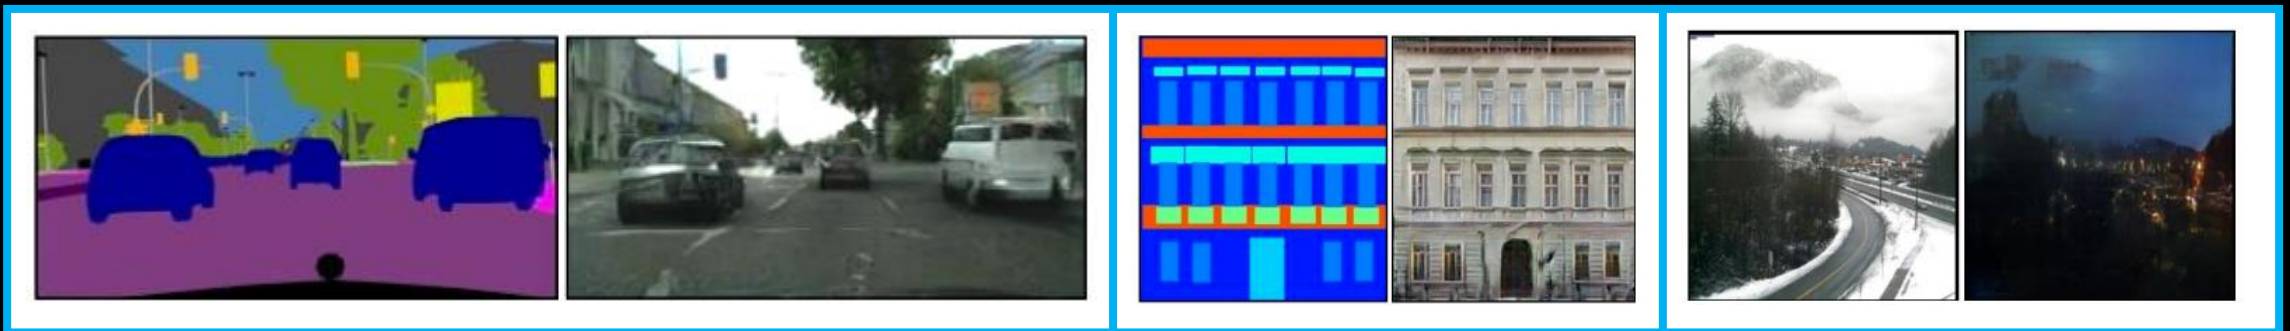
(PS: U-Net is from 2015, this [paper](#) from 2018)



Pix2Pix

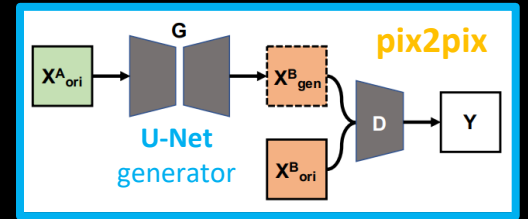


- General purpose **image-to-image translation**:
 - From their paper: “Many problems in image processing, computer graphics, and computer vision can be posed as “*translating*” an input image into a *corresponding output image*.”
 - Translating without specifically defining the rules – data-driven translating between two domains by showing **paired** examples:
 - [Image from A, Corresponding Image from B] * N samples




(PS: there is also an un-paired version called **CycleGAN** = images from two domains, but they don't have to correspond)

Pix2Pix



- General purpose **image-to-image translation**:



Pix2Pix

- **Online demos:** affinelayer.com/pixsrv/
- **Colab notebooks:**
 - Training and using Pix2Pix on Colab: our repo / [Demo1_pix2pix-keras-v2.ipynb](#)
- **Papers:** [pix2pix](#), [pix2pixHD](#) (with high. res.), [vid2vid](#) (with frame to frame consistency)

Style Transfer

- Previously we saw, how a Deep Convolutional network separates where it saves *high-level* and *low-level* representations (this is due to the *Conv->Pool* combo – each convolutional layer serves as an image-filter).

Style Transfer

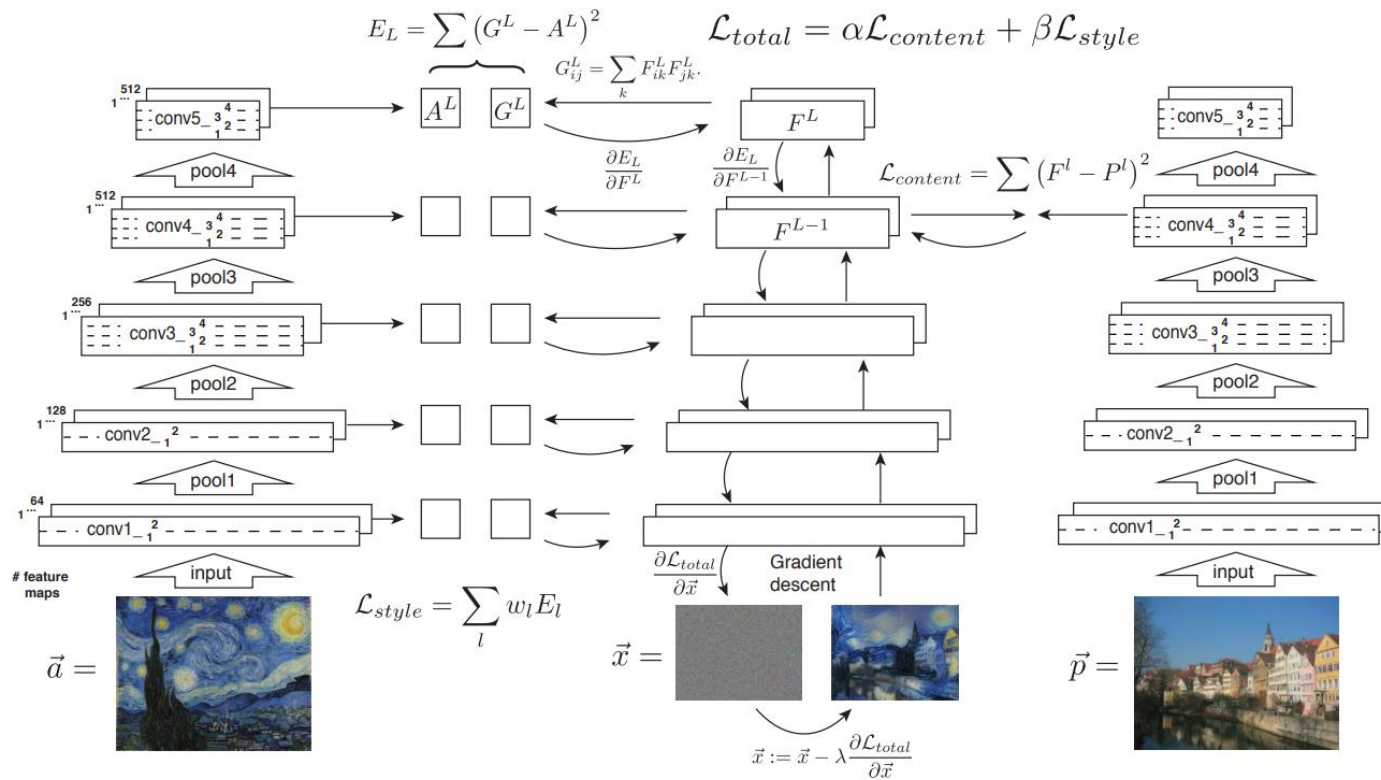
- Previously we saw, how a Deep Convolutional network separates where it saves *high-level* and *low-level* representations (this is due to the *Conv->Pool* combo – each convolutional layer serves as an image-filter).
 - **High-level information** – contents of the whole image – **image content**
 - Encode one image extracting its *content information*, the feature responses in deeper layers of the network

Style Transfer

- Previously we saw, how a Deep Convolutional network separates where it saves *high-level* and *low-level* representations (this is due to the *Conv->Pool* combo – each convolutional layer serves as an image-filter).
 - **High-level information** – contents of the whole image – **image content**
 - Encode one image extracting its *content information*, the feature responses in deeper layers of the network
 - **Low-level information** – details used inside the image, texture – **image style**
 - Encode another image extracting the *style information*, feature response alongside a selection of layers

(PS: One possible similarity if the content of low and high frequencies in music converted to spectrograms)

Style Transfer



Style Transfer

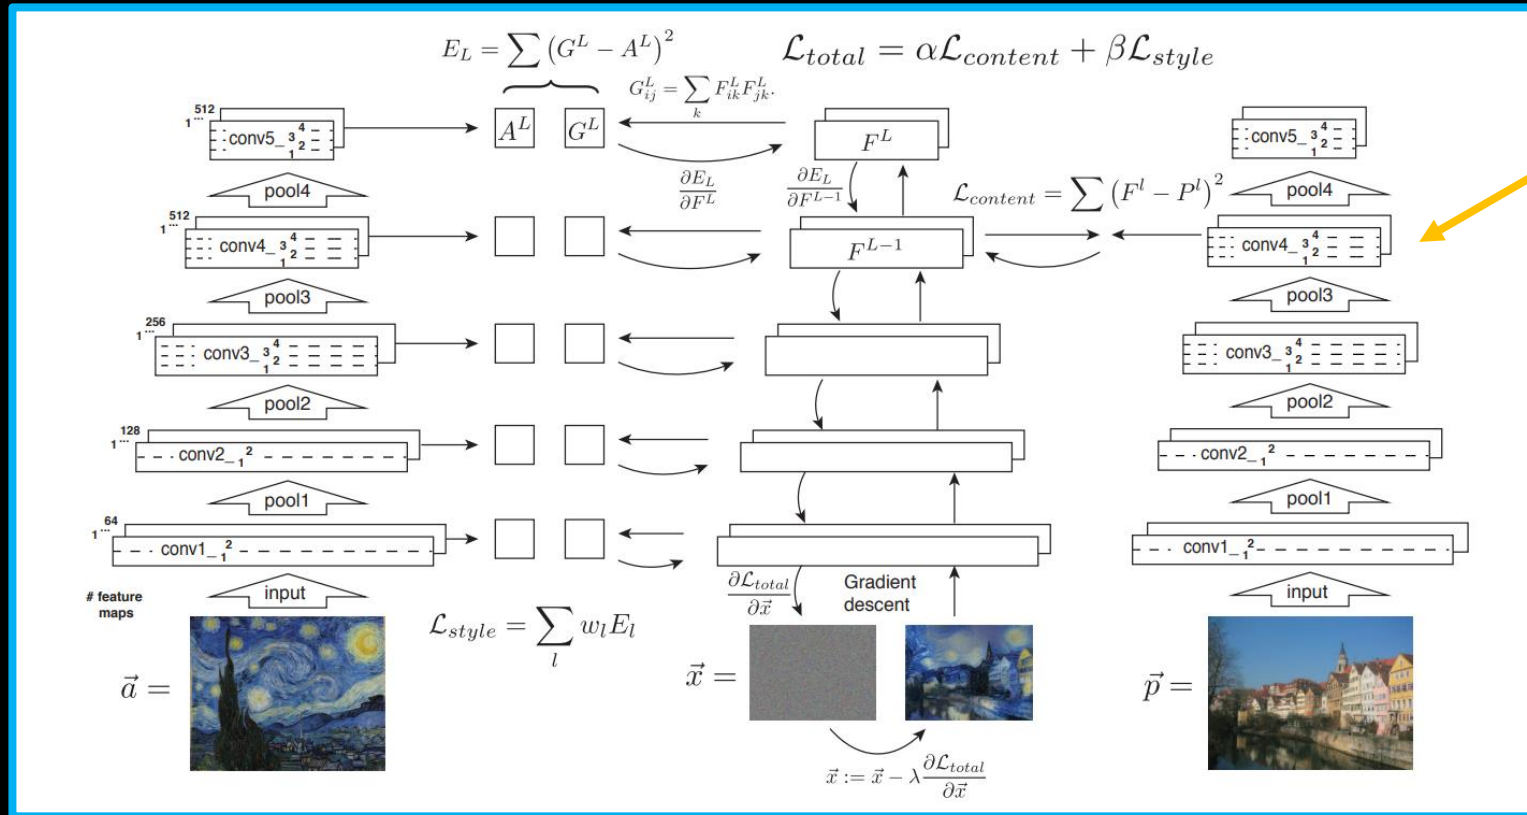


Image content =
the feature
responses in deeper
layers of the VGG
network

Style Transfer

Image style =
feature response
alongside a
selection of layers

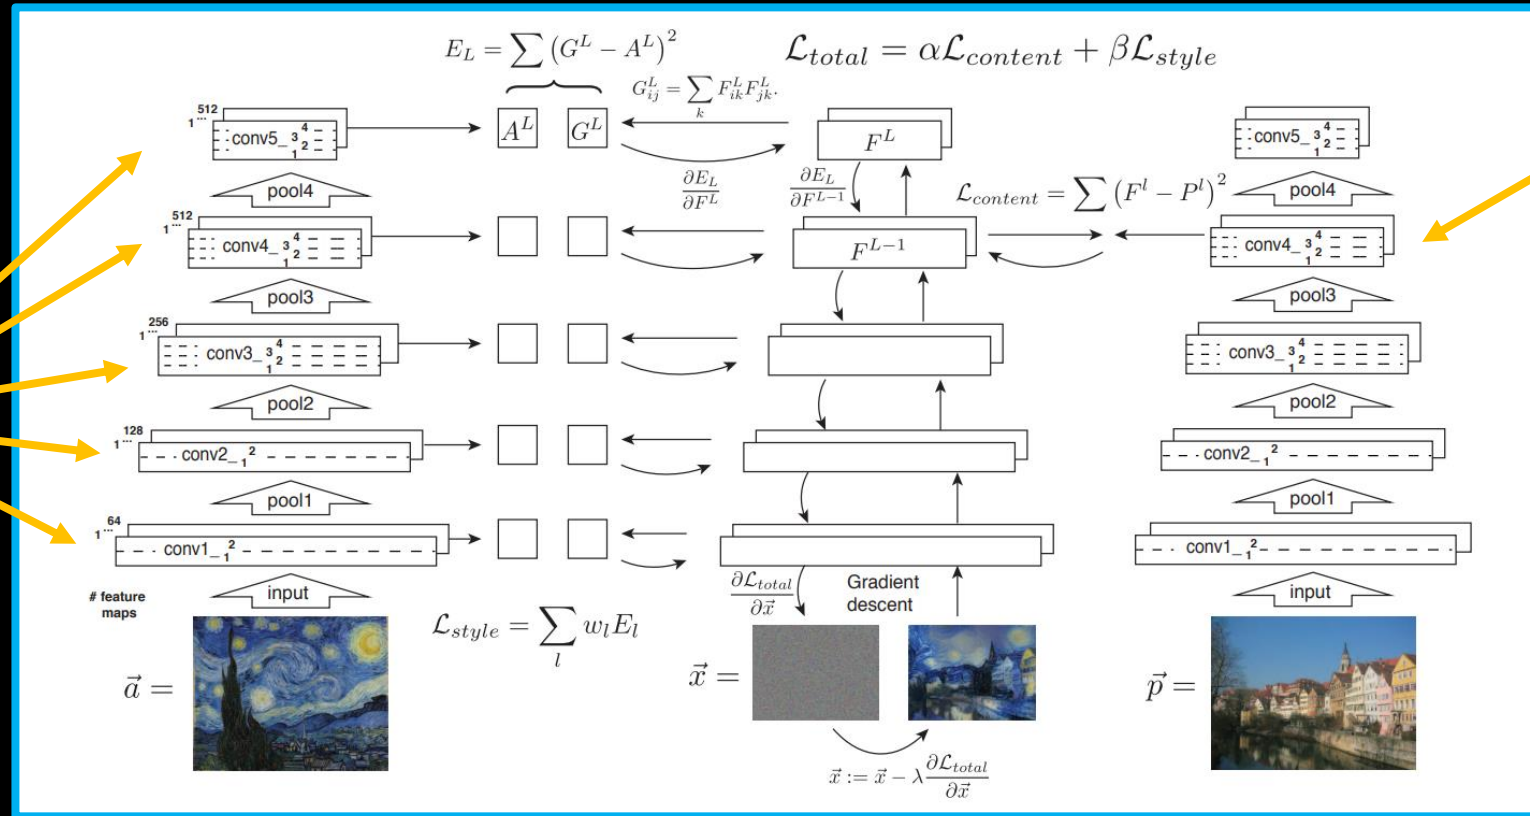


Image content =
the feature
responses in deeper
layers of the VGG
network

Style Transfer

Image style =
feature response
alongside a
selection of layers

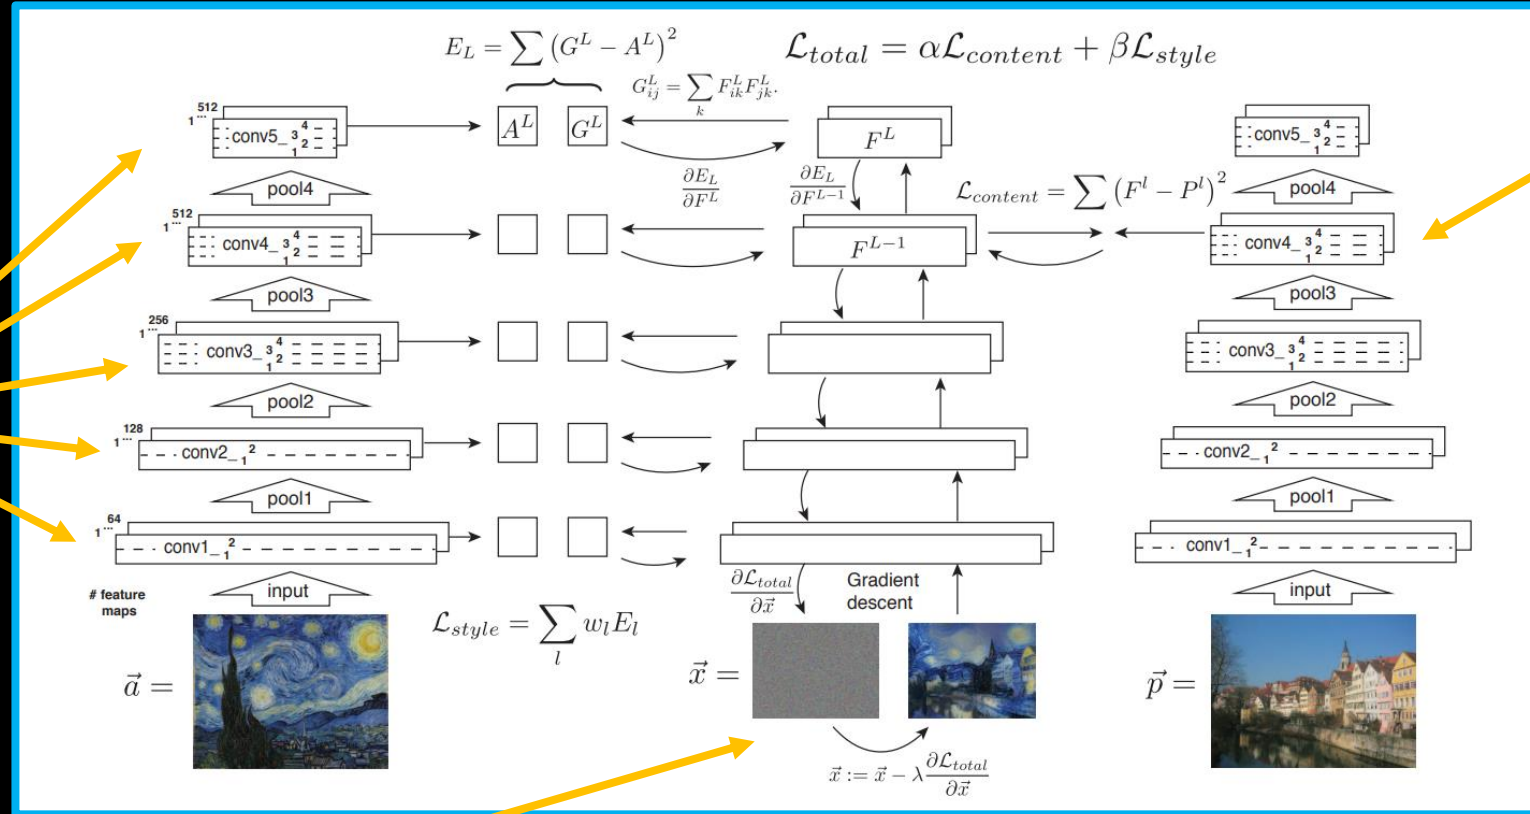


Image content =
the feature
responses in deeper
layers of the VGG
network

Optimizing random noise image to have the same responses as those features we saved.

- Iteratively we will create a new image which has the style responses similar to our encoded style features + content responses similar to our encoded content features

Style Transfer

Image style =
feature response
alongside a
selection of layers

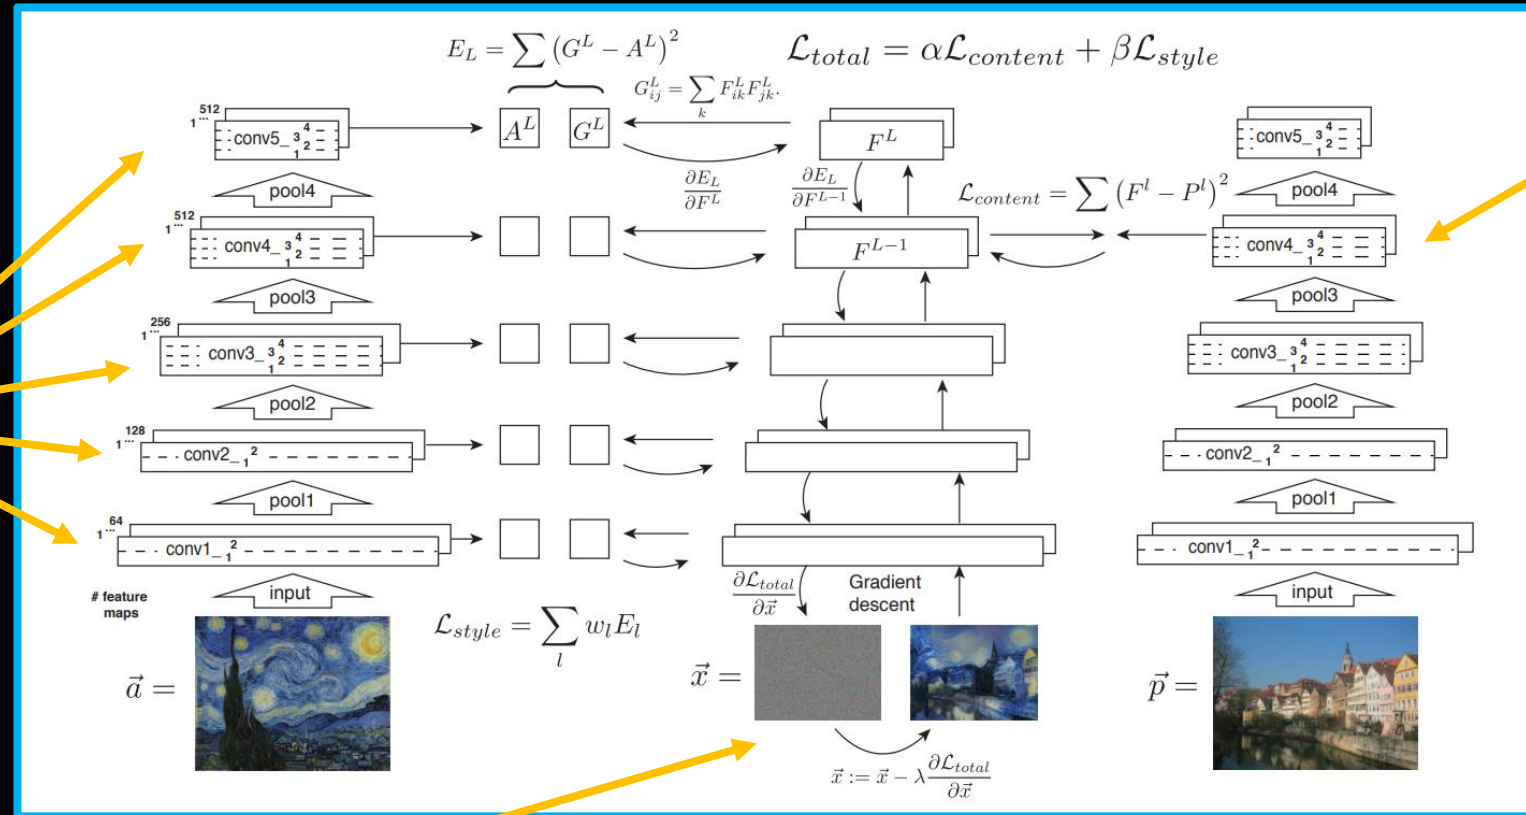


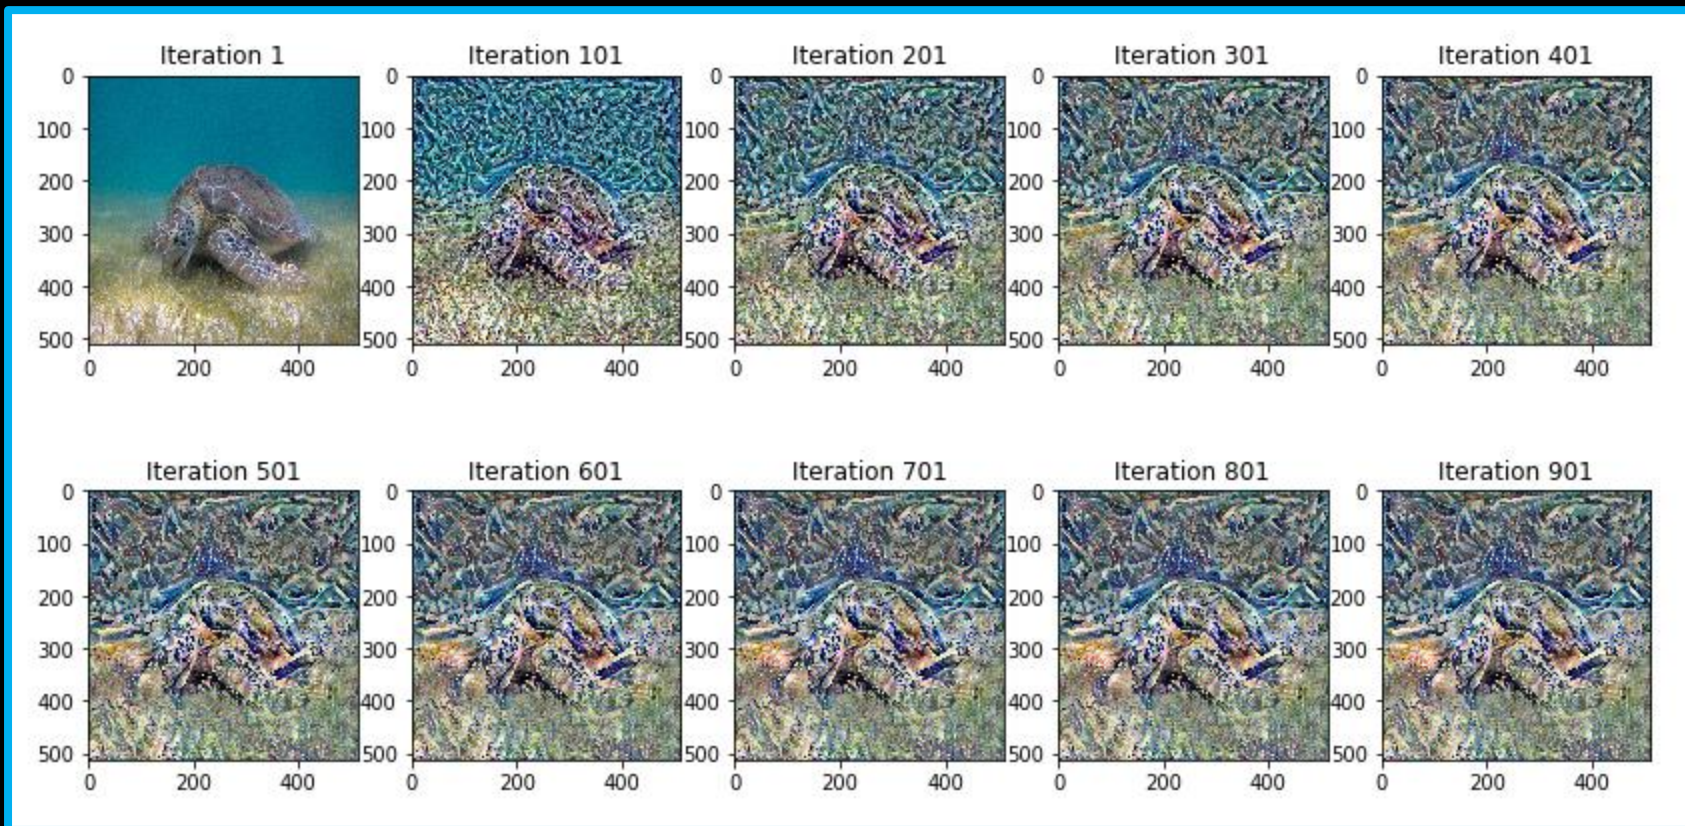
Image content =
the feature
responses in deeper
layers of the VGG
network

Optimizing random noise image to have the same responses as those features we saved.

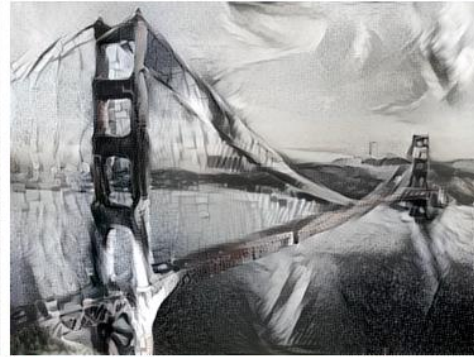
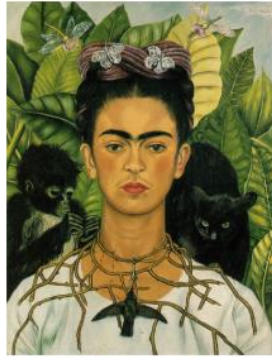
- Iteratively we will create a new image which has the style responses similar to our encoded style features + content responses similar to our encoded content features

This is **relatively slow** (iterative optimization of the input image), later papers sped it up by using an **image2image translation method** with feed-forward networks.

Style Transfer



Style Transfer




Style Transfer

- **Online demos:** deepart.io/latest/
- **Colab notebooks:**
 - Basic style transfer (with arbitrary images): ArtML / [style transfer keras.ipynb](#)
 - Fast style transfer (with pretrained styles): ArtML / [fast-style-transfer](#)
- **Papers:** [style transfer \(2015\)](#), [fast style transfer \(2016\)](#)


Deep Dream

- Deep dream is a method which was originally used to **visualize what a network has learned**. It works on **image optimization** principle (as did the first version of Style Transfer).

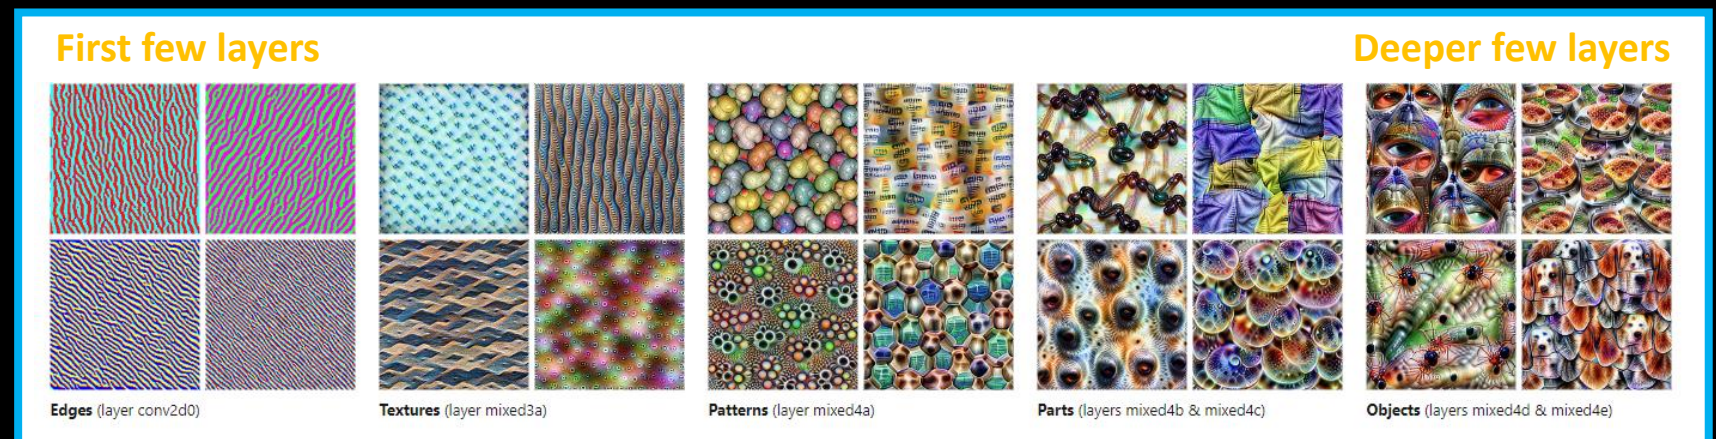


Deep Dream

- Deep dream is a method which was originally used to **visualize what a network has learned**. It works on **image optimization** principle (as did the first version of Style Transfer).




Using **Convolutional network** (GoogLeNet) trained for classification on ImageNet:



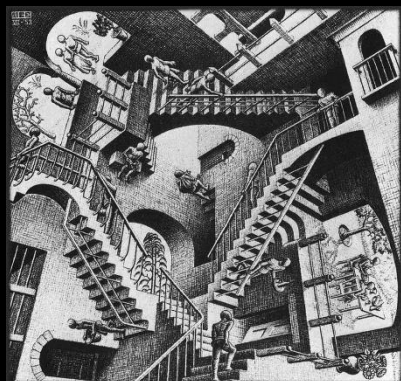
(PS: image [source](#), features [viz.](#), illustrational [video](#))

Deep Dream

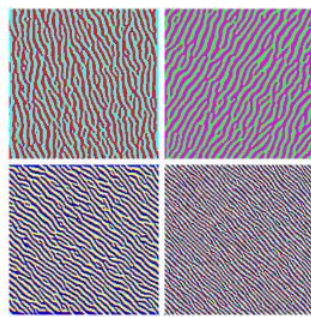
- Deep dream is a method which was originally used to **visualize what a network has learned**. It works on **image optimization** principle (as did the first version of Style Transfer).



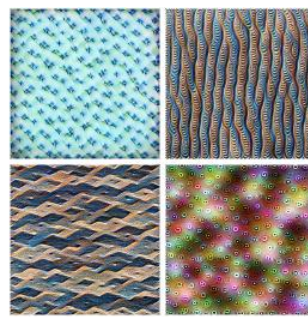
Change this **original image** so that it activates a **selected feature** with the highest possible force!



First few layers



Edges (layer conv2d0)



Textures (layer mixed3a)



Patterns (layer mixed4a)



Parts (layers mixed4b & mixed4c)

Deeper few layers

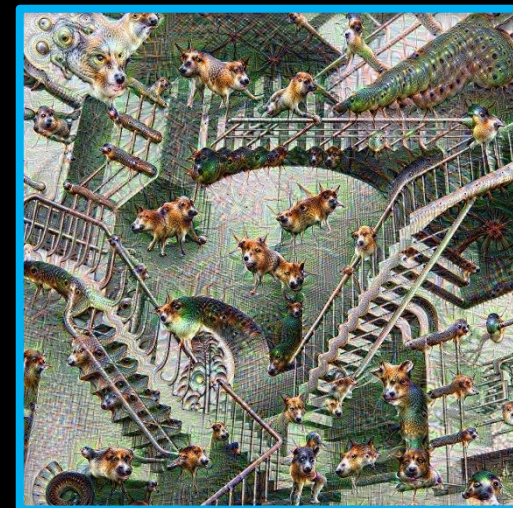


Objects (layers mixed4d & mixed4e)

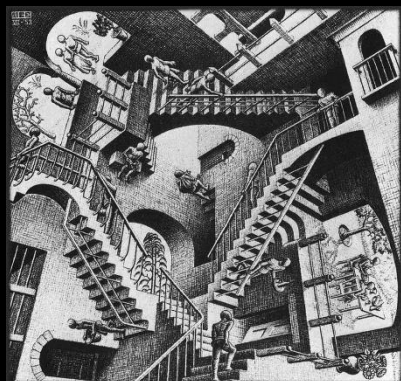
(PS: image [source](#), features [viz.](#), illustrational [video](#))

Deep Dream

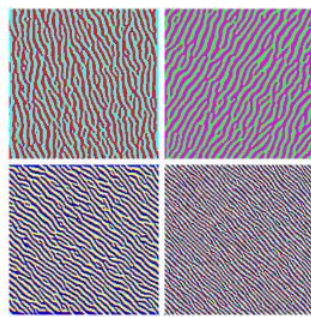
- Deep dream is a method which was originally used to **visualize what a network has learned**. It works on **image optimization** principle (as did the first version of Style Transfer).



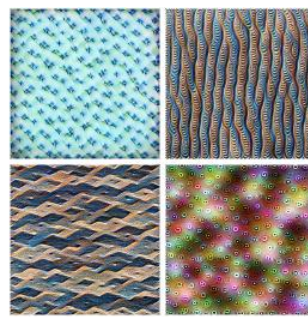
Change this **original image** so that it activates a **selected feature** with the highest possible force!



First few layers



Edges (layer conv2d0)



Textures (layer mixed3a)



Patterns (layer mixed4a)



Parts (layers mixed4b & mixed4c)

Deeper few layers



Objects (layers mixed4d & mixed4e)



selected feature

(PS: image [source](#), features [viz.](#), illustrational [video](#))

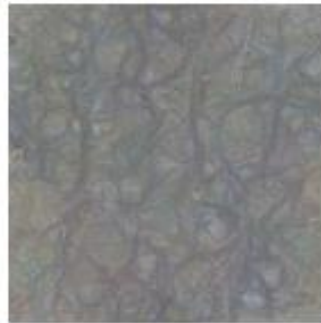
Deep Dream

Iterations:

Starting from random noise, we optimize an image to activate a particular neuron (layer mixed4a, unit 11).



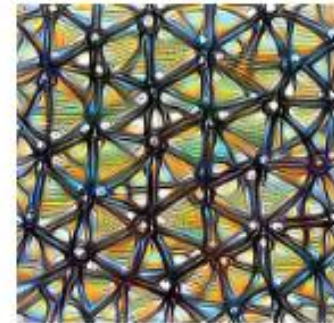
Step 0



Step 4



Step 48



Step 2048

Deep Dream

As a network visualization technique ...

Deep dreaming can reveal information stored inside the network ...
Recall shapes and content information from the originally used dataset.

- Model released by Yahoo to identify NSFW content (without releasing the sensitive information about the dataset): github.com/yahoo/open_nsfw
- Which was (*of course* ...) soon followed by using a deep-dream–like technique: open_nsfw.gitlab.io/
 - Interesting usage of the network to generate *suggestive imagery* (changing shapes of landscapes imagery, etc.)



Deep Dream

- **Online demos:** dreamscopeapp.com *(not 100% sure if it's not a style transfer of deep dream like effect)*
- **Colab notebooks:**
 - Deep dream a photo: [as a ML4A guide](#)
 - Alternative code: ArtML / [neural-synth-clustering-v2.ipynb](#)
- **Reading:** distill.pub/2017/feature-visualization/

Exploring Machine Intelligence

Week 6, Generative Models II



Detailed usage of Progressive GAN


Practicum: Generative Models II.

This week's focus will be in learning how to use **Progressive Growing GAN** with all practical steps included:


- data processing: [1 Dataset Processing.ipynb](#)
- model training: [2 Train ProgressiveGAN.ipynb](#)
- model inference and interpolations: [3 ProgressiveGAN Inference.ipynb](#)

Progressively Growing GAN:

Video [link?t=82](#)



Latent space



If we **found images** with some property and without this property **in the latent space** (for example: smiling / neutral expression) ...

PS: (important) difference with AutoEncoders is that **we can't encode** them into the latent space

We can find clusters and check their relative positions:

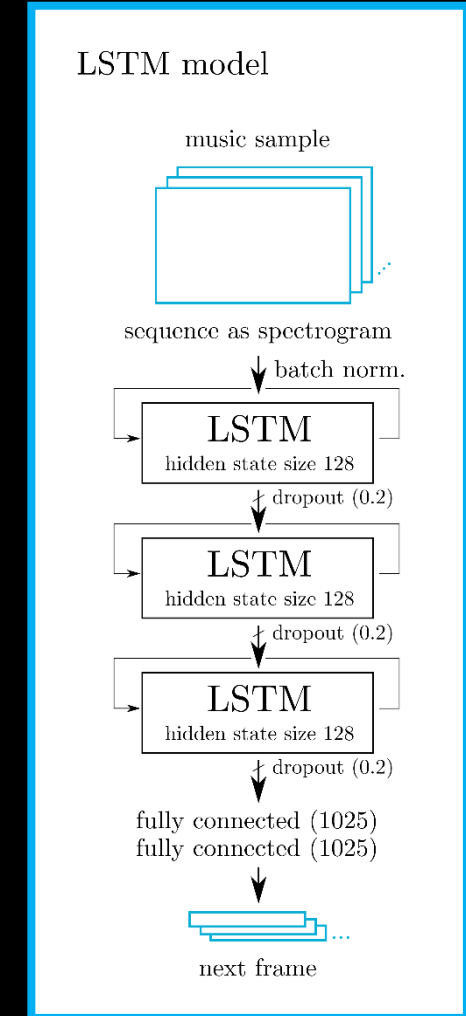
$$v = \text{centroid of } c_1 - \text{centroid of } c_2$$

- We can call this **latent space arithmetic**

Next class

More generative models:

- Sequential modelling



Links and additional readings:

- **Bonus readings:**

- **Feature Visualization**, Distill – [blog](#)
- **Sensory Optimization**: Neural Networks as a Model for Understanding and Creating Art – [paper](#)
- About **Pix2Pix** on ML4A – [blog with code](#)

The end